Microsoft

Microsoft Partner Project Ready

Data Integration and Transformation with Data Factory in Microsoft Fabric

Day 01 of 02

Technofocus

Course Plan and Learning Objectives

Day 1

240 mins

Module 1 - Introduction to Data Factory in Microsoft Fabric

- Microsoft Fabric - The unified data platform for AI transformation
- Data Factory in Microsoft Fabric

Module 2 - Ingest data with Data Factory in Microsoft Fabric

- Data ingestion with Pipelines
- Copy activity with Pipelines
- Use parameters and expressions in Pipelines
- Data ingestion with Copy Job
- Mirroring databases in Microsoft Fabric
- Execute, monitor and troubleshoot Pipelines
- Data Pipeline storage event triggers
- REST API capabilities and CI/CD for Pipelines
- Fabric Pipelines vs. ADF/Synapse Pipelines

Hands-on labs

- Use case 01: Implementing Medallion Architecture with Data Factory in Microsoft Fabric for scalable data processing

Day 2

240 mins

Module 3 - Data Transformation with Dataflows Gen2

- Dataflows Gen2 in Microsoft Fabric
- Fast Copy in Dataflows Gen2
- Dataflow Gen2 with CI/CD and Git integration support
- Monitor your Dataflows
- Copilot for Data Factory in Microsoft Fabric
- Workflow Orchestration with Apache Airflow job

Module 4 - Migrate to Data Factory in Microsoft Fabric

- Plan your migration from Azure Data Factory to Data Factory in Microsoft Fabric
- Migrate from Dataflow Gen1 to Dataflow Gen2
- Move queries from Dataflow Gen1 to Dataflow Gen2
- Ingest data into Microsoft Fabric using the Azure Data Factory Copy Activity

Hands-on labs

- Use Case 02: Data Factory solution for moving and transforming data with dataflows and data pipelines

04
Migrate to Data Factory in Microsoft Fabric

Plan your migration from Azure Data Factory to Data Factory in Microsoft Fabric

Migrate from Azure Data Factory to Data Factory in Microsoft Fabric

Design Principle

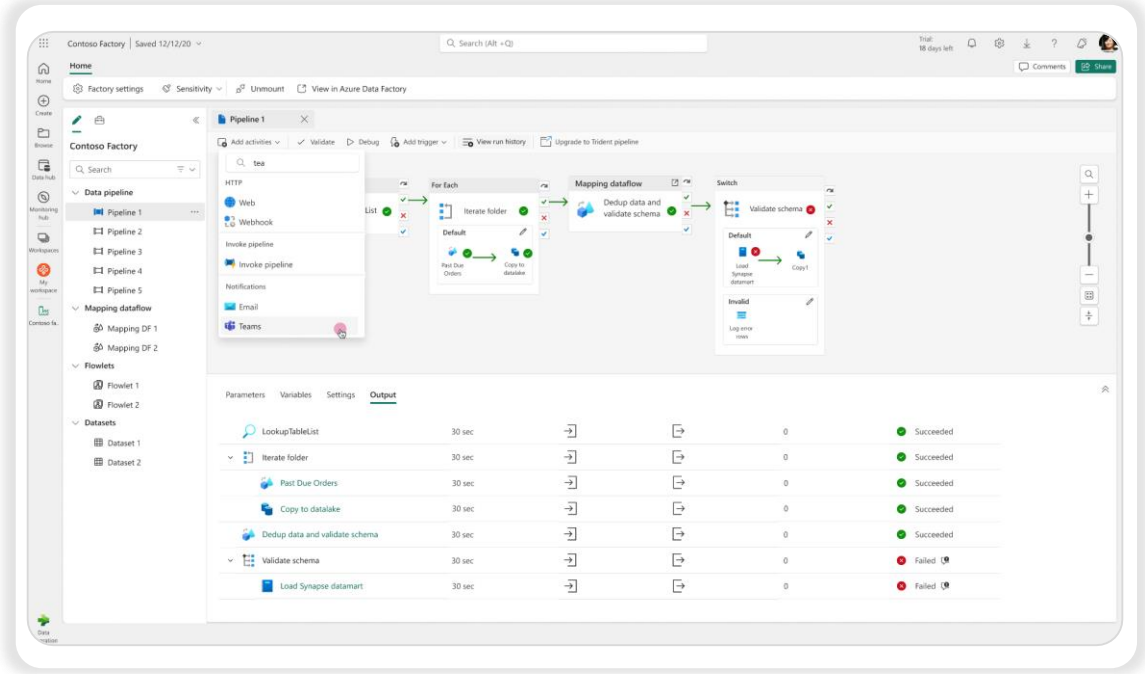

Frictionless Transition

Provides a seamless integration of Azure Data Factory into Microsoft Fabric, allowing users to easily manage their data pipelines.

Modernized Experience

Provides a modernized user interface and pipeline management experience for the pipelines available within the mounted Azure data factory

New Capabilities

Provides with exciting opportunities to streamline their data integration processes and unblock new capabilities.

Migration benefits

- Easy routing of messages during pipeline execution with email and Teams activities
- Built-in continuous integration and delivery (CI/CD) features (deployment pipelines) don't require external integration with Git repositories
- Workspace integration with your OneLake data lake enables single-pane-of-glass easy analytics management
- Refreshing your semantic data models is easy in Fabric with a fully integrated pipeline activity

License requirement

- Fabric pipelines and Dataflow Gen2 require a Microsoft Fabric premium capacity workspace

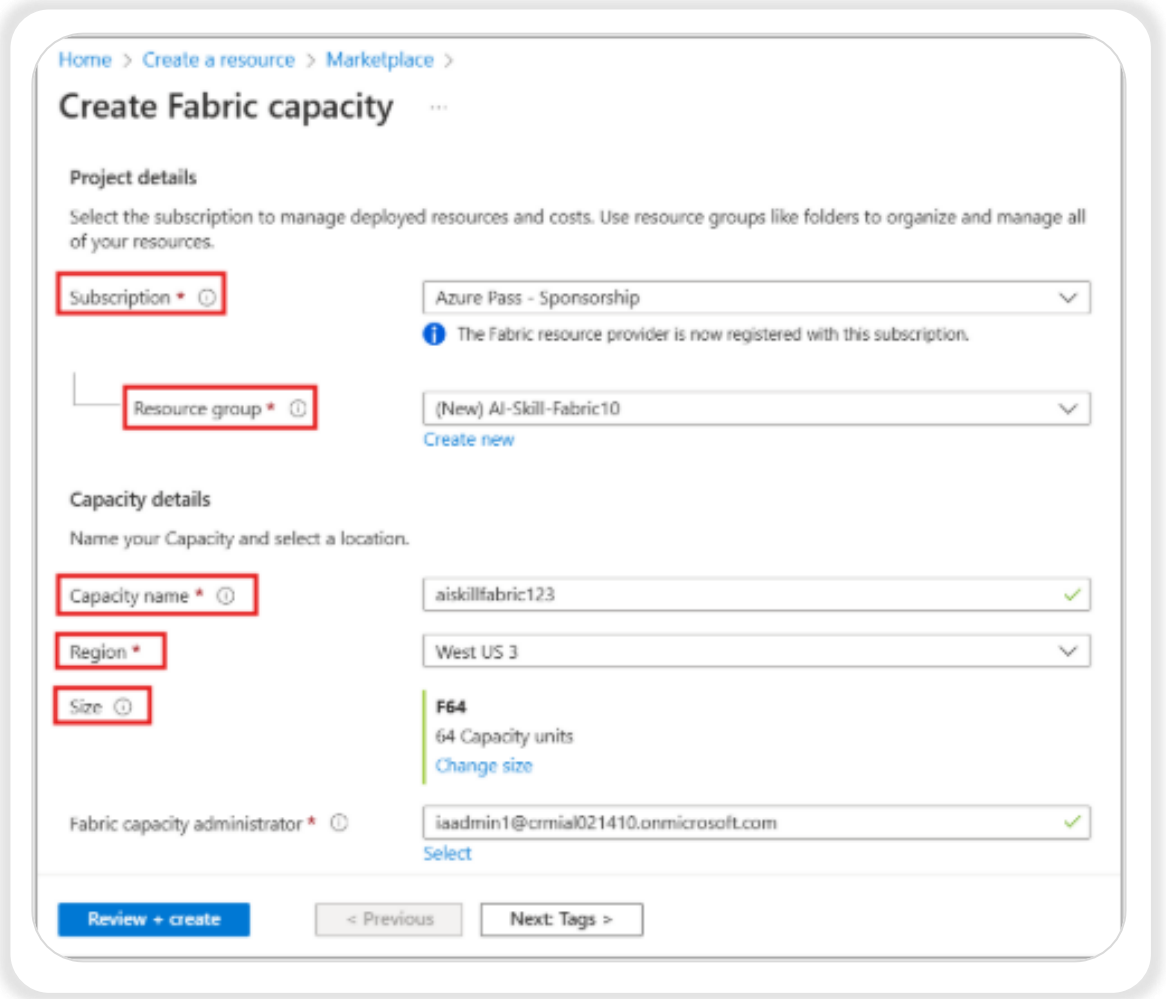

Platform differences between ADF and Data Factory in Fabric

Features	Azure Data Factory	Data Factory in Fabric
	PaaS product	SaaS product
Integration Runtimes	IRs are configuration objects that represent compute	Default is to use cloud-based compute
Pipelines	Fundamental component for the primary workflow and orchestration of your ADF processes	Pipelines in Fabric Data Factory are nearly identical to ADF. The JSON definition of pipelines differs slightly
Linked Services	Linked Services define the connectivity properties needed to connect to your data stores	Recreate definitions as Connections
Datasets	Datasets define the shape, location, and contents of your data in ADF	Don't exist as entities in Fabric
Dataflows	Data Flows are built on the Synapse Spark infrastructure	Power Query based Dataflows

Platform differences between ADF and Data Factory in Fabric

Features	Azure Data Factory	Data Factory in Fabric
Triggers	Features are similar in Fabric although the underlying implementation is different	Triggers only exist as a pipeline concept. Data Activator framework is used
Debugging	Separate debug mode you find in ADF pipelines and data flows	Debugging pipelines is simpler in Fabric than in ADF
Change Data Capture	Preview feature, move data quickly in an incremental manner	To migrate CDC artifacts to Fabric Data Factory, recreate the artifacts as Copy job
Azure Synapse Link	Not available in ADF	Recreate the Azure Synapse Link artifacts as Mirroring items
SQL Server Integration Services (SSIS)	Lift-and-shift your SSIS packages into the cloud using the ADF SSIS IR	Don't have the concept of IRs, so this functionality isn't possible today
Invoke pipeline activity	A common activity: Execute pipeline activity	Enhanced activity: Invoke pipeline activity

Migration scenario – Migrate ADF pipelines and data flows

- Modernize your ETL environment from the ADF factory PaaS model to the new Fabric SaaS model
- The primary factory items to migrate: pipelines and data flows
- Other fundamental factory elements to migrate: linked services, integration runtimes, datasets, and triggers

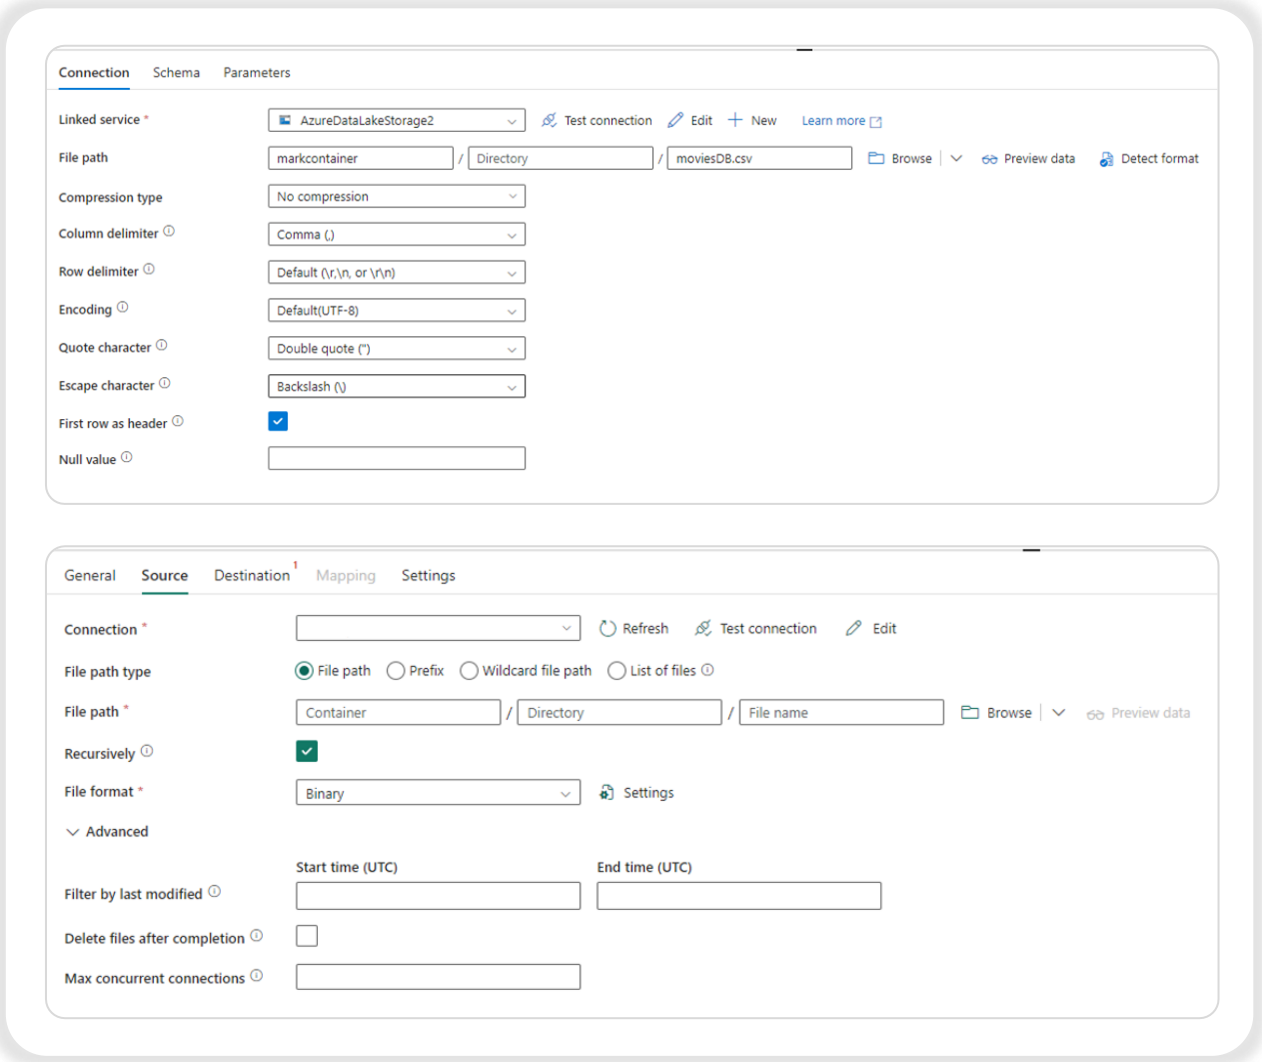

Migration scenario - Migrate ADF with CDC, SSIS, and Airflow

- These features serves different data integration needs, but require special attention when migrating
- CDC is a top-level ADF concept but in Fabric, you see this capability as the Copy job
- Airflow is available in both

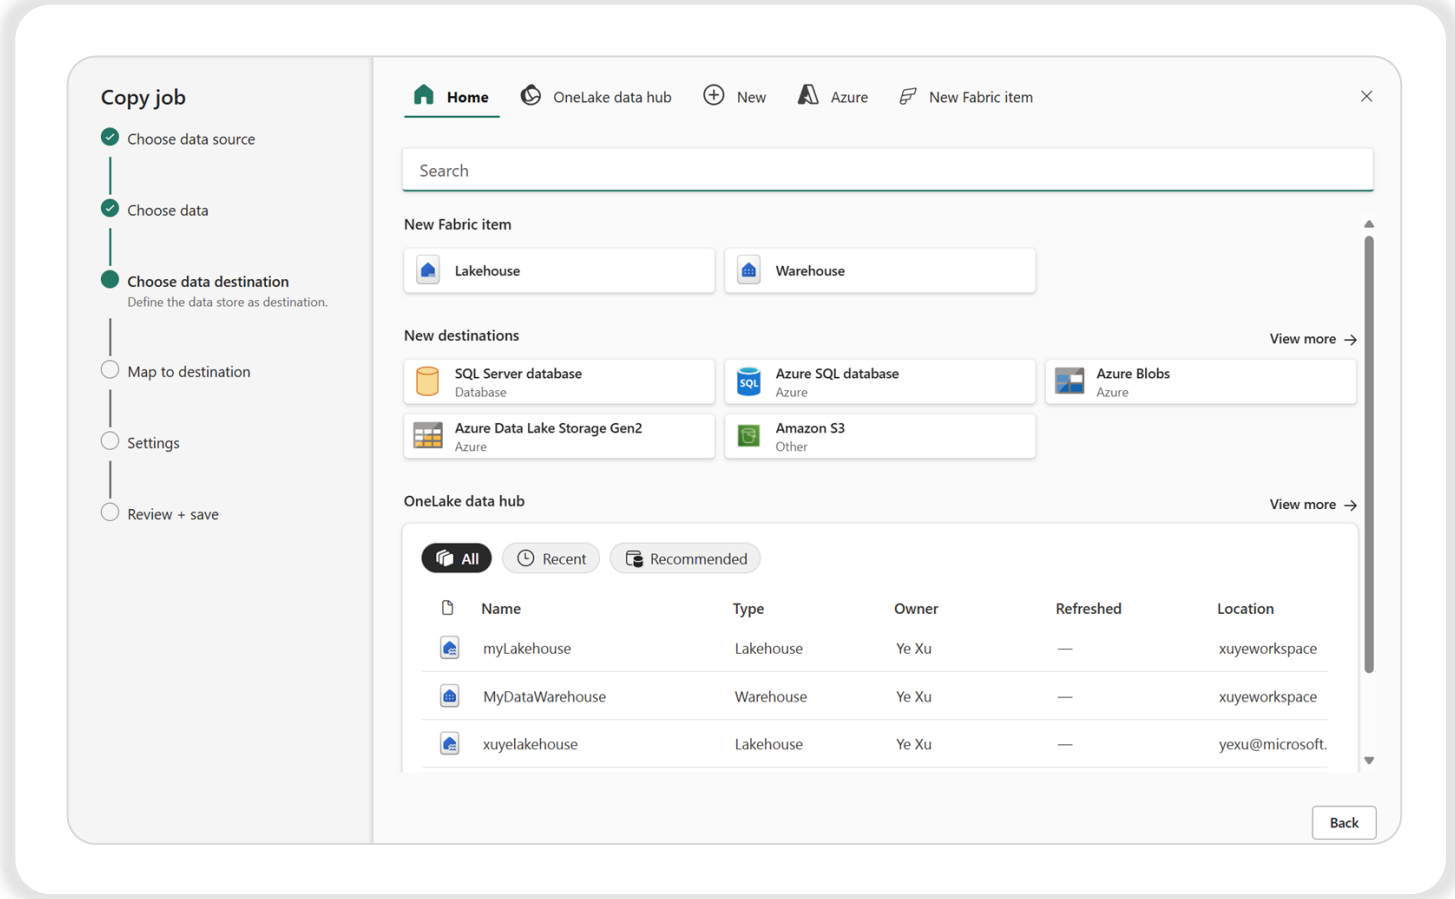

Migration scenario – Git-enabled Data Factory migration to Fabric

- Fabric provides two primary ways to enable CI/CD, both at the workspace level: Git integration and built-in deployment pipelines
- Existing Git repo from ADF doesn't work with Fabric

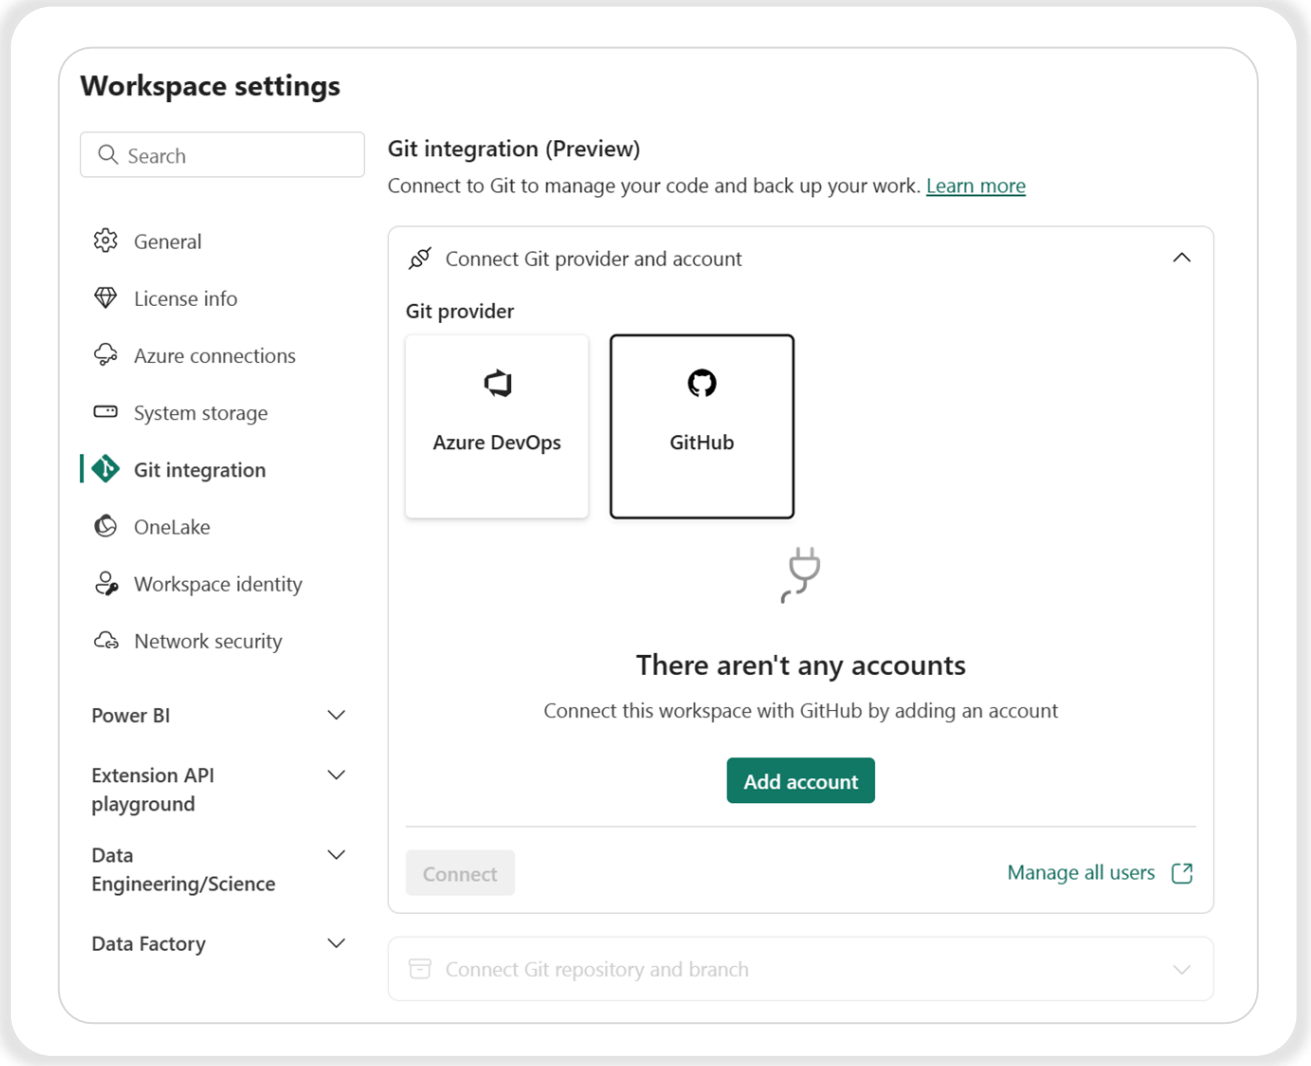

Activity continuity

- Data Factory in Fabric aims to offer comprehensive capabilities while maintaining compatibility with ADF
- New activities in Fabric Data Factory
  - Outlook
  - Teams
  - Semantic model refresh
  - Dataflow Gen2

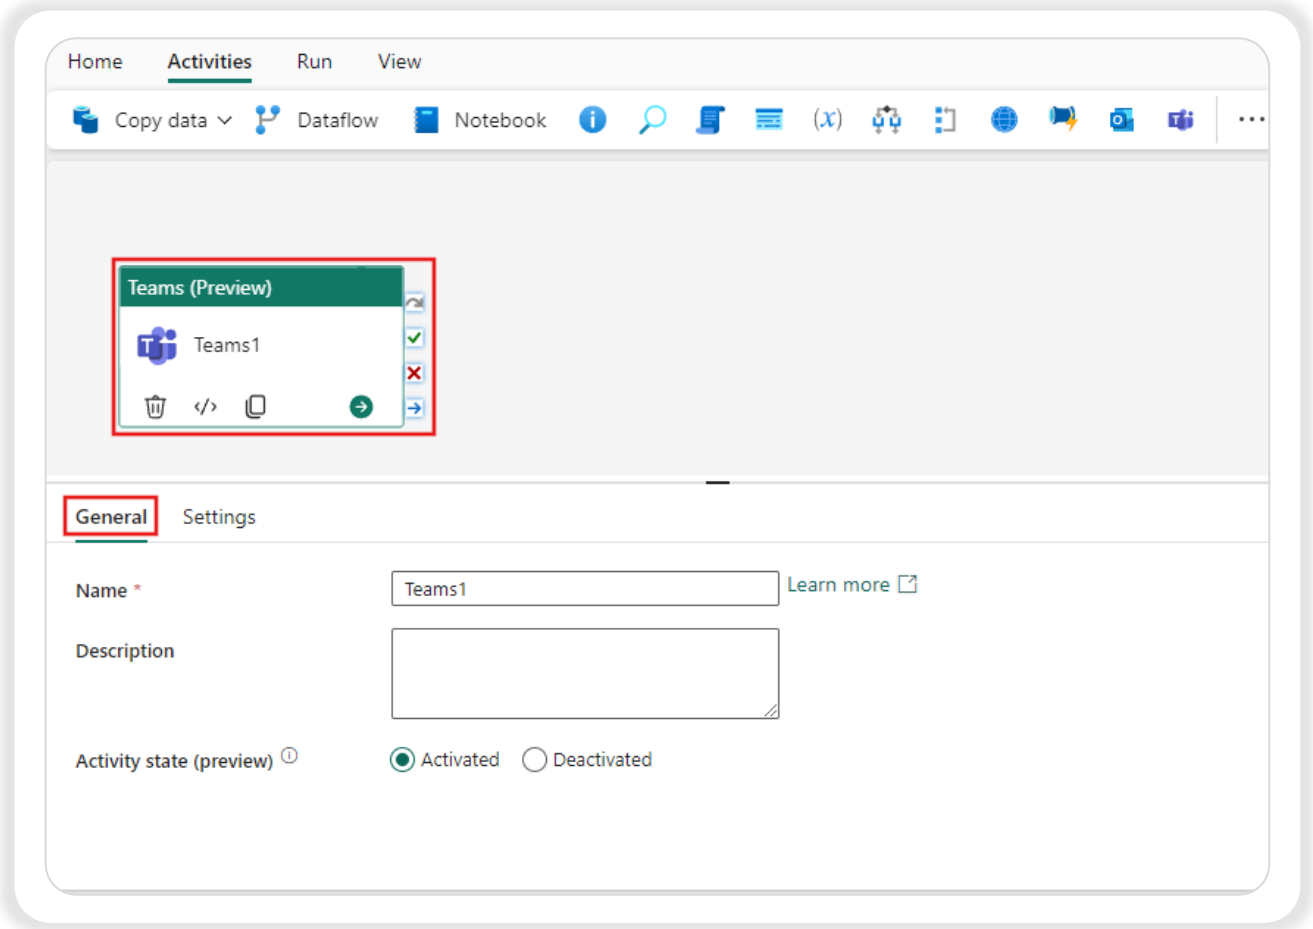

Connector continuity

With Data Factory in Microsoft Fabric, data pipelines provide connectivity to a rich set of data sources

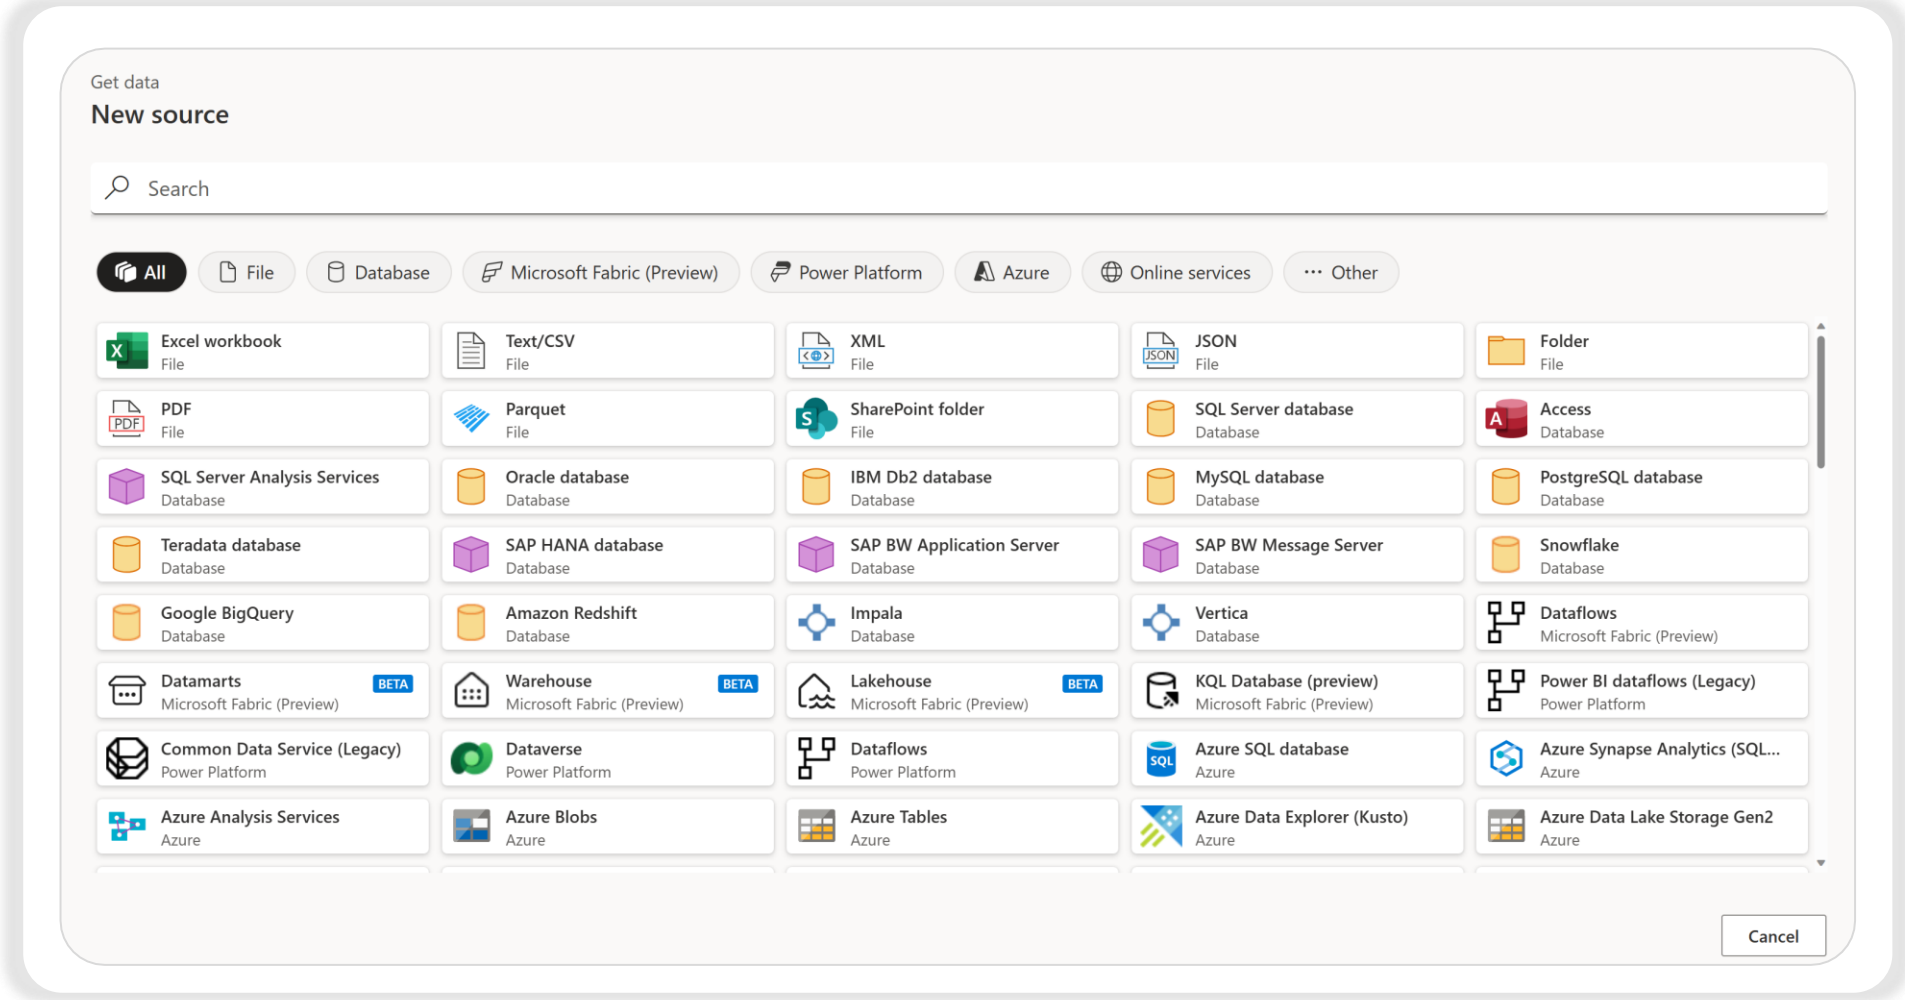

Mapping Data Flow transformations to the Dataflow Gen2

- When first starting to author Dataflows, you can also use the Global search box
- Mapping data flow transformations supported in Dataflow Gen2
  - Multiple inputs/outputs
  - Schema modifier
  - Formatters
  - Row modifier
  - Flowlets
  - Destination

Mapping Data Flow	Dataflow Gen2
New branch	Reference
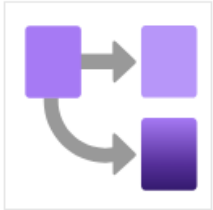	 **Important:** Right select a query and select Reference.
Aggregate	**Group by**
	

Mapping data flows to PySpark Notebooks

The Mapping data flow to Microsoft Fabric notebook is a command line tool to convert Mapping data flow script code to Spark Scala and PySpark code

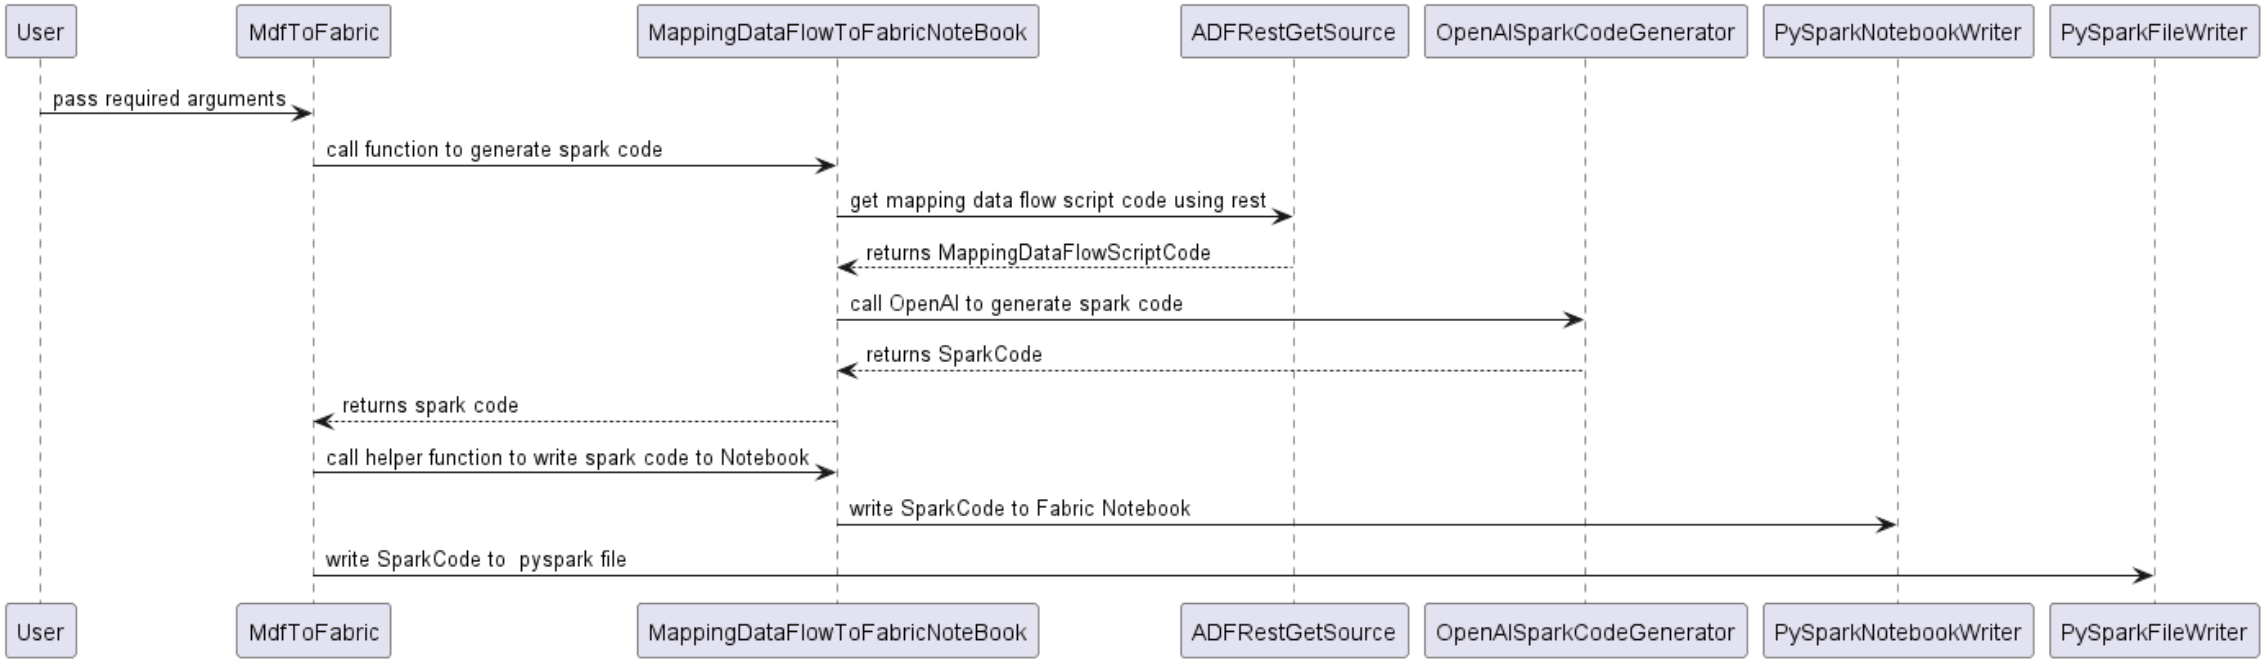

Mounting

Within Fabric, you will have the ability to mount existing Synapse Pipeline/Azure Data Factory artifact. This allows you to still preserve the feature in ADF but also leverage new features in Microsoft Fabric.

- The mount provides the interface UI in Fabric to allow customers to develop/edit their existing ADF investment and it is NOT migration.
- The mount provide the interface UI in Microsoft Fabric platform to allow customer to edit
- Development can happen within Fabric or within the existing ADF workspace.
- User is allowed to add activity that is only available in Microsoft Fabric – E-mail notification, Team Activity, destination output, etc
- Everything user do right now, whether custom Dev/Ops, meta driven pipeline should work with mount, since this is nothing for than a UI, but still using the Azure platform.
- The backend is still using the Azure platform, until it is migrated to the Microsoft Fabric platform
- ADF customer can also invoke activity in Microsoft Fabric.

Use Case - Mount/Unmount

This feature allows users to select Azure Data Factory as their preferred mount point, providing them with an easy and seamless way to use their ADF in Fabric

Start from Fabric Data Integration homepage

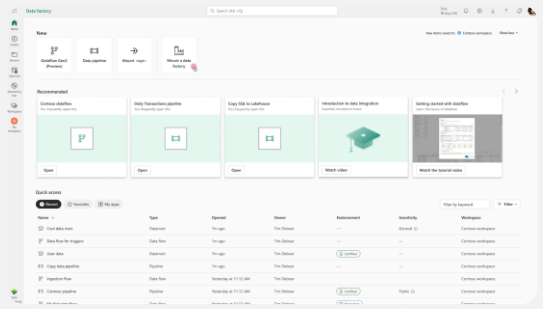

Choose target data factories

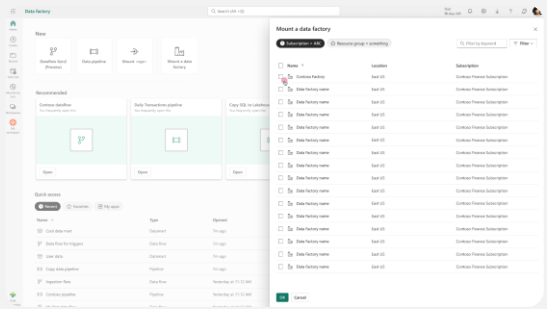

Mount to the Fabric workspace

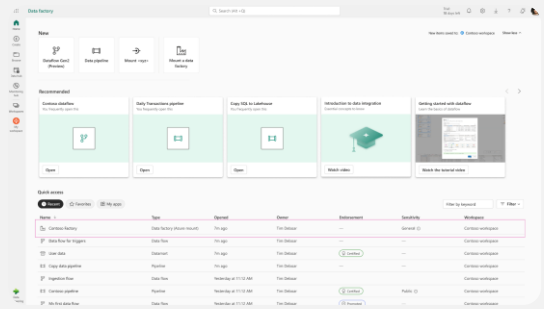

- The mount can be filter at the subscription, resource group. It allows user to mount into the Fabric WS with just one click.
- User also has the option to unmount to remove from the Fabric WS
- This mount option will also be available in existing ADF/Synapse Pipeline to the Fabric WS

Use Case - Frictionless Transition & Hybrid Development

This feature provides a seamless integration of Azure Data Factory into the Fabric platform, allowing users to easily manage their data pipelines and work with their data in a centralized location.

Open a mounted data factory with object explore view

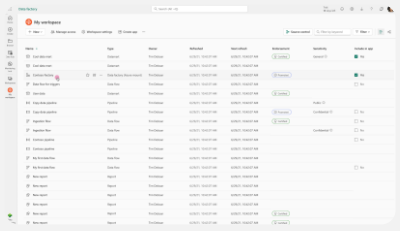

Select a pipeline view on the canvas

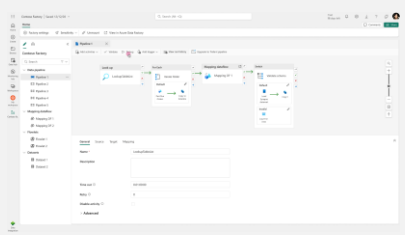

Run a mounted pipeline

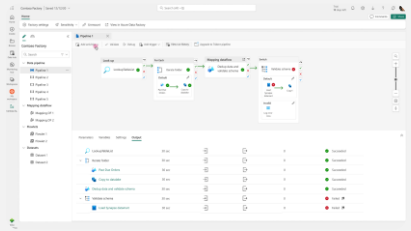

Add a Teams activity

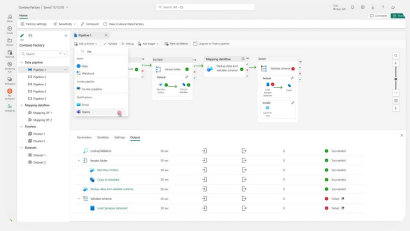

- User can select a mounted ADF in the Fabric WS, all the elements of ADF is visible
- All the changes made in Fabric will be sync back in ADF
- Customer can run the pipeline in Trident allowing them to experience both platforms at the same time

Use Case - Central Monitoring

This feature provides a central place for monitoring of all pipeline run without having to switch between different platforms.

View mounted data factory pipeline in monitoring hub

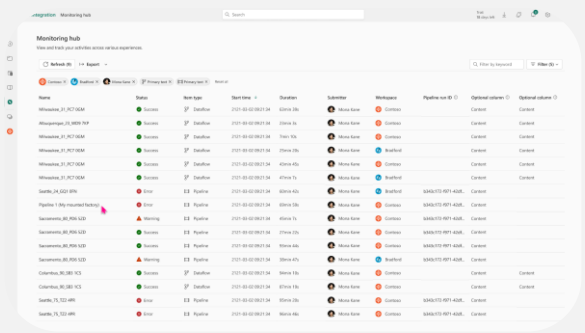

View mapping dataflow execution details in monitoring page

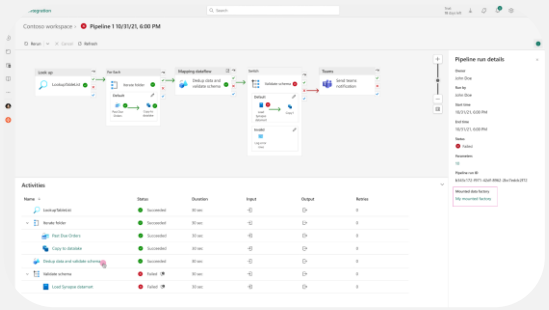

View mounted data factory pipeline execution details in pipeline monitoring page

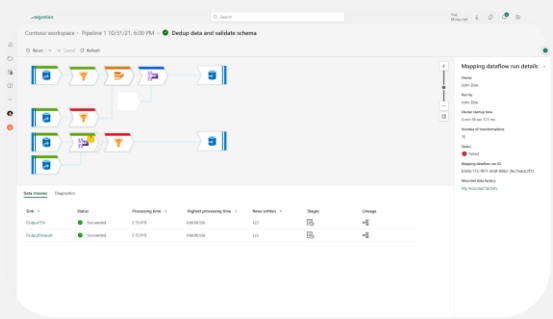

Demo

Mount your existing ADF instances directly to a Fabric workspace

https://daily.fabric.microsoft.com/home?trident=1&experience=data-factory

Data Factory Home

Search

Trial: 23 days left

16

Home

Workspaces

OneLake data hub

Monitor

Real-Time hub

DemoWorkspace01

Data Factory

Welcome to Microsoft Fabric

Get started with Data Factory

Create a workspace with a predesigned template called a task flow. Task flows keep your items organized. - ...

Build your own task flow

General

Basic data analytics

Data analytics using a SQL ...

Medallion

Ev

Recommended items to create

Data pipeline

Azure Data Factory (preview)

Apache Airflow job (Preview)

Current workspace: DemoWorkspace01

Items will be saved to this workspace.

Show less

Learn more about Data Factory

Learn more about Data Factory

Learn to use Data Factory

Get started with Data Factory

Create your first dataflow

Get started with dataflows

Create your first data pipeline

Get started with data pipelines

Learn to monitor Data Factory

Monitoring your Data Factory

Migrate from Dataflow Gen1 to Dataflow Gen2

Feature overview

Feature	Dataflow Gen2	Dataflow Gen1
Author dataflows with Power Query	✓	✓
Shorter authoring flow	✓	
AutoSave and background publishing	✓	
Data destinations	✓	
Improved monitoring and refresh history	✓	
Integration with data pipelines	✓	
High-scale compute	✓	
Get Data via Dataflows connector	✓	✓
Direct Query via Dataflows connector		✓
Incremental refresh	✓	✓
AI Insights support		✓

Migration guidance

- Compile an inventory of your dataflows and dependent items
  - Dataflows as a source in Power BI
  - Dataflows as a source in Power Apps
  - Dataflows as a source in Excel
- Consider using Power Query templates

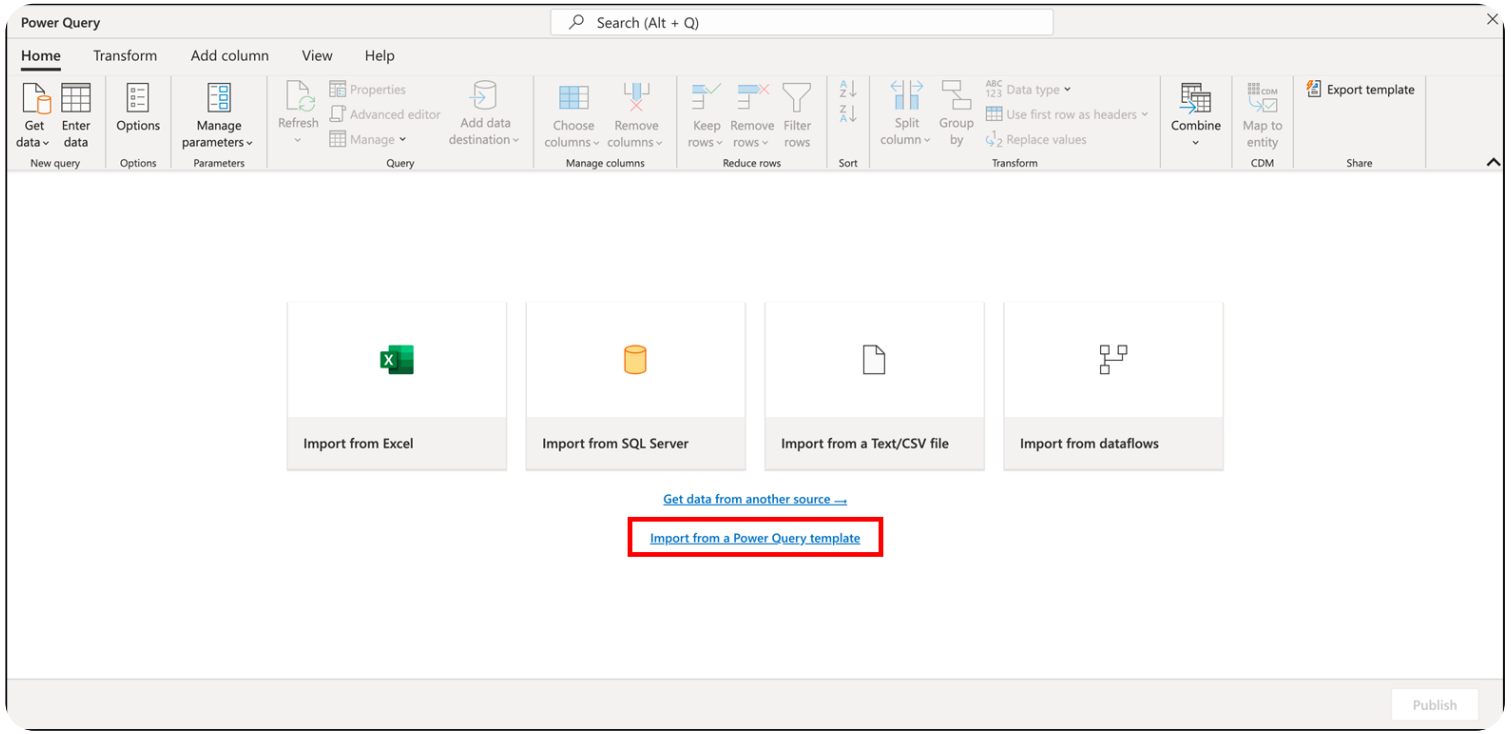

Migrate from Dataflow Gen1 to Dataflow Gen2: Migration scenarios

- **Personal or team usage**

  Use dataflows to automate data ingestion and preparation tasks, allowing them to focus on data analysis and insights

- **Departmental usage**

  Use dataflows to manage larger data sources and complex transformations

- **Enterprise usage**

  Use dataflows in ingesting vast amounts of data across multiple departments at scale

Migration scenario - Personal or team usage

- **Problem Statement**
  - The dataflow creators want to take advantage of the advanced capabilities of Dataflow Gen2 for authoring purposes
  - At the same time, they plan to temporarily continue using dataflow tables as a data source during a phased migration
- **Solution**
  - Update the workspace ID, if a new workspace is created to store the new dataflow
  - Update existing solutions from the original (Gen1) dataflow ID to the new (Gen2) dataflow ID

Power Query M Copy

```
let
    Source = PowerPlatform.Dataflows(null),
    Workspaces = Source{[Id="Workspaces"]}[Data],
    Workspace = Workspaces{[workspaceId="<enter new workspace ID>"]}[Data],
    DataflowId = Workspace{[dataflowId="<enter new dataflow ID"]}[Data],
    DimDateTable = DataflowId{[entity="DimDate", version=""]}[Data]
in
    DimDateTable
```

Migration scenario - Departmental usage

- **Problem Statement**
  - The dataflow creators want to take advantage of the advanced capabilities of Dataflow Gen2 for authoring
  - Efficiently sharing and outputting the dataflow tables to a Fabric lakehouse
  - This method takes advantage of OneLake shortcuts
- **Solution**
  - Replace linked tables with OneLake shortcuts, which provide downstream consumers with direct access to the data
  - Update existing solutions and transition queries by replacing the PowerPlatform.Dataflows or PowerBI.Dataflows functions with the Lakehouse.Contents data access function in Fabric

Power Query M

Copy

```
let
  Source = Lakehouse.Contents([]),
  WorkspaceId = Source{[workspaceId="<0000aaaa-11bb-cccc-dd22-eeeeee333333>"]}[Data],
  LakehouseId = WorkspaceId{[lakehouseId="1111bbbb-22cc-dddd-ee33-ffffff444444"]}[Data],
  DimCustomerTable = LakehouseId{[Id="DimCustomer", ItemKind="Table"]}[Data]
in
  DimCustomerTable
```

Migration scenario - Enterprise usage

- **Problem Statement**
  - The dataflow creators want to take advantage of the advanced capabilities of Dataflow Gen2 for authoring
  - Outputting and sharing dataflow tables from a Fabric warehouse that has granular user permissions
- **Solution**
  - Grant data access through the SQL compute engine's granular permissions, which provide more selective access to certain users by restricting access to specific tables and schemas, as well as implementing RLS and CLS
  - Update existing solutions and transition queries by replacing the PowerPlatform.Dataflows or PowerBI.Dataflows function with the Fabric.Warehouse data access function in Fabric

Power Query M

Copy

```
let
  Source = Fabric.Warehouse([]),
  WorkspaceId = Source{[workspaceId="0000aaaa-11bb-cccc-dd22-eeeeee333333"]}[Data],
  WarehouseId = WorkspaceId{[warehouseId="1111bbbb-22cc-dddd-ee33-ffffff444444"]}[Data],
  DimCustomerTable = WarehouseId{[Schema="dbo", Item="DimCustomer"]}[Data]
in
  DimCustomerTable
```

Migrate to Dataflow Gen2 (CI/CD) using Save As (Preview)

- Save a Dataflow Gen2 or Gen2 (CI/CD) as a new Dataflow Gen2 (CI/CD)
- Save a Dataflow Gen1 as a new Dataflow Gen2 (CI/CD)

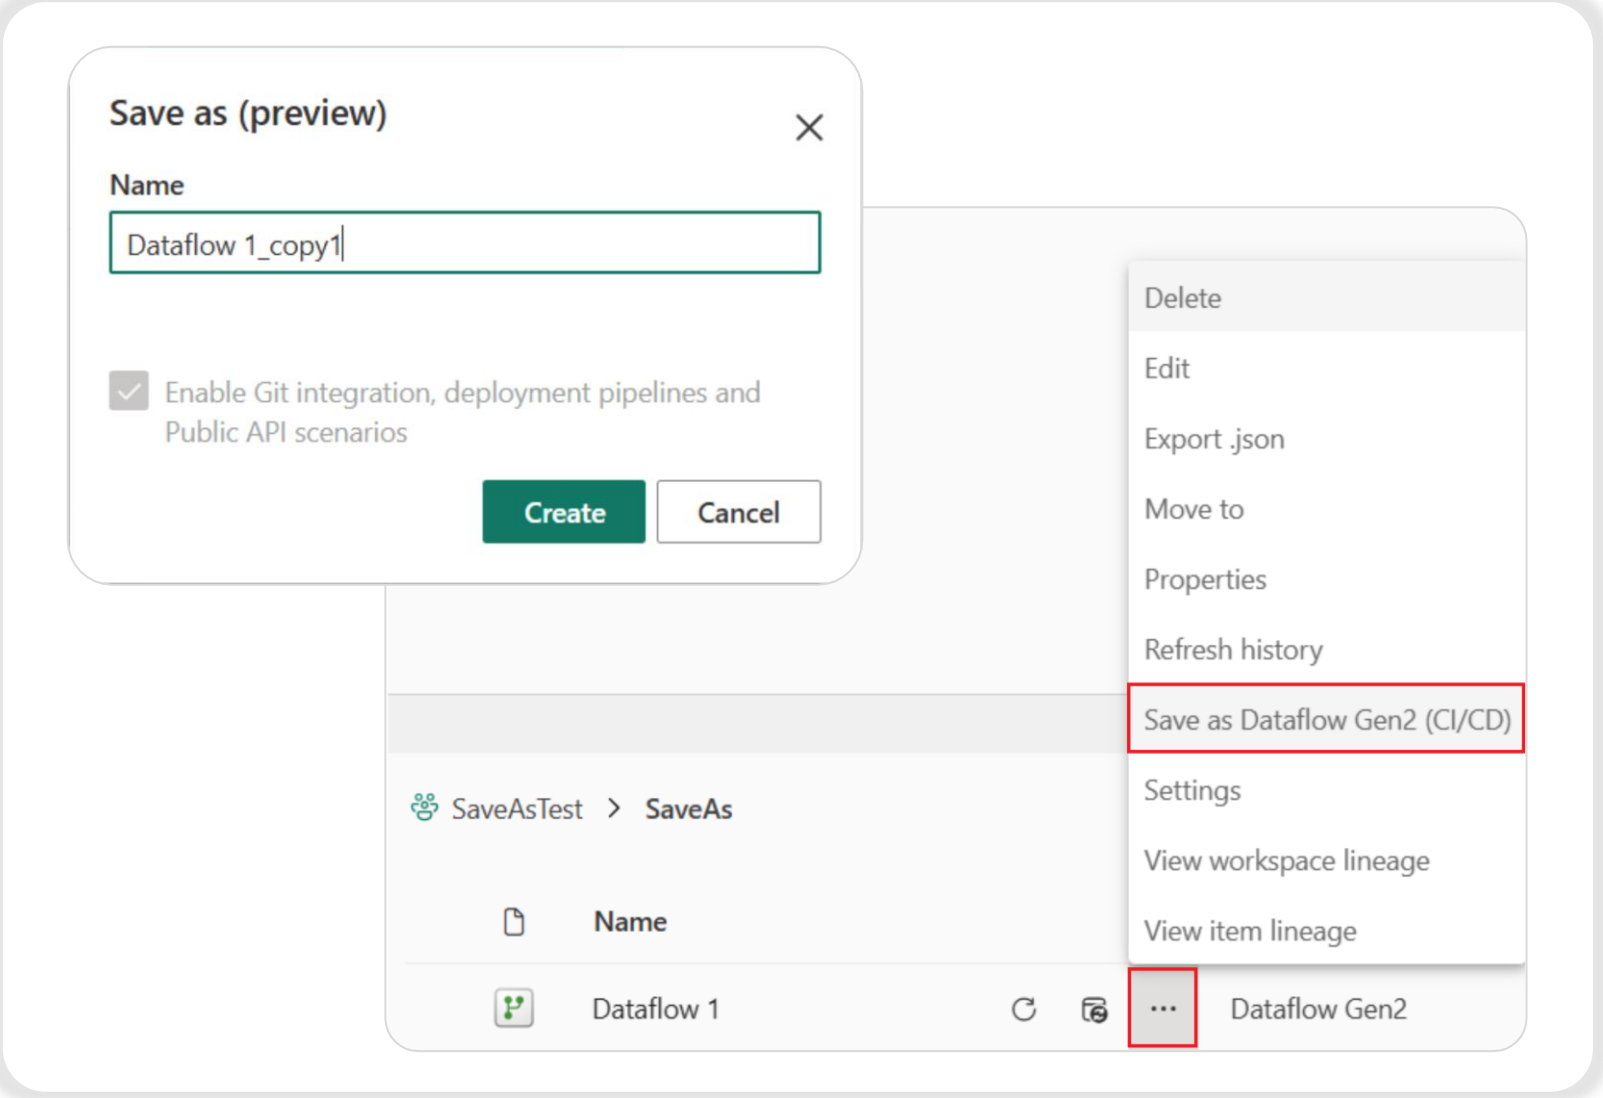

Move queries from Dataflow Gen1 to Dataflow Gen2

Move queries from Dataflow Gen1 to Dataflow Gen2

- Use the export template feature
- Copy and paste existing Dataflow Gen1 queries

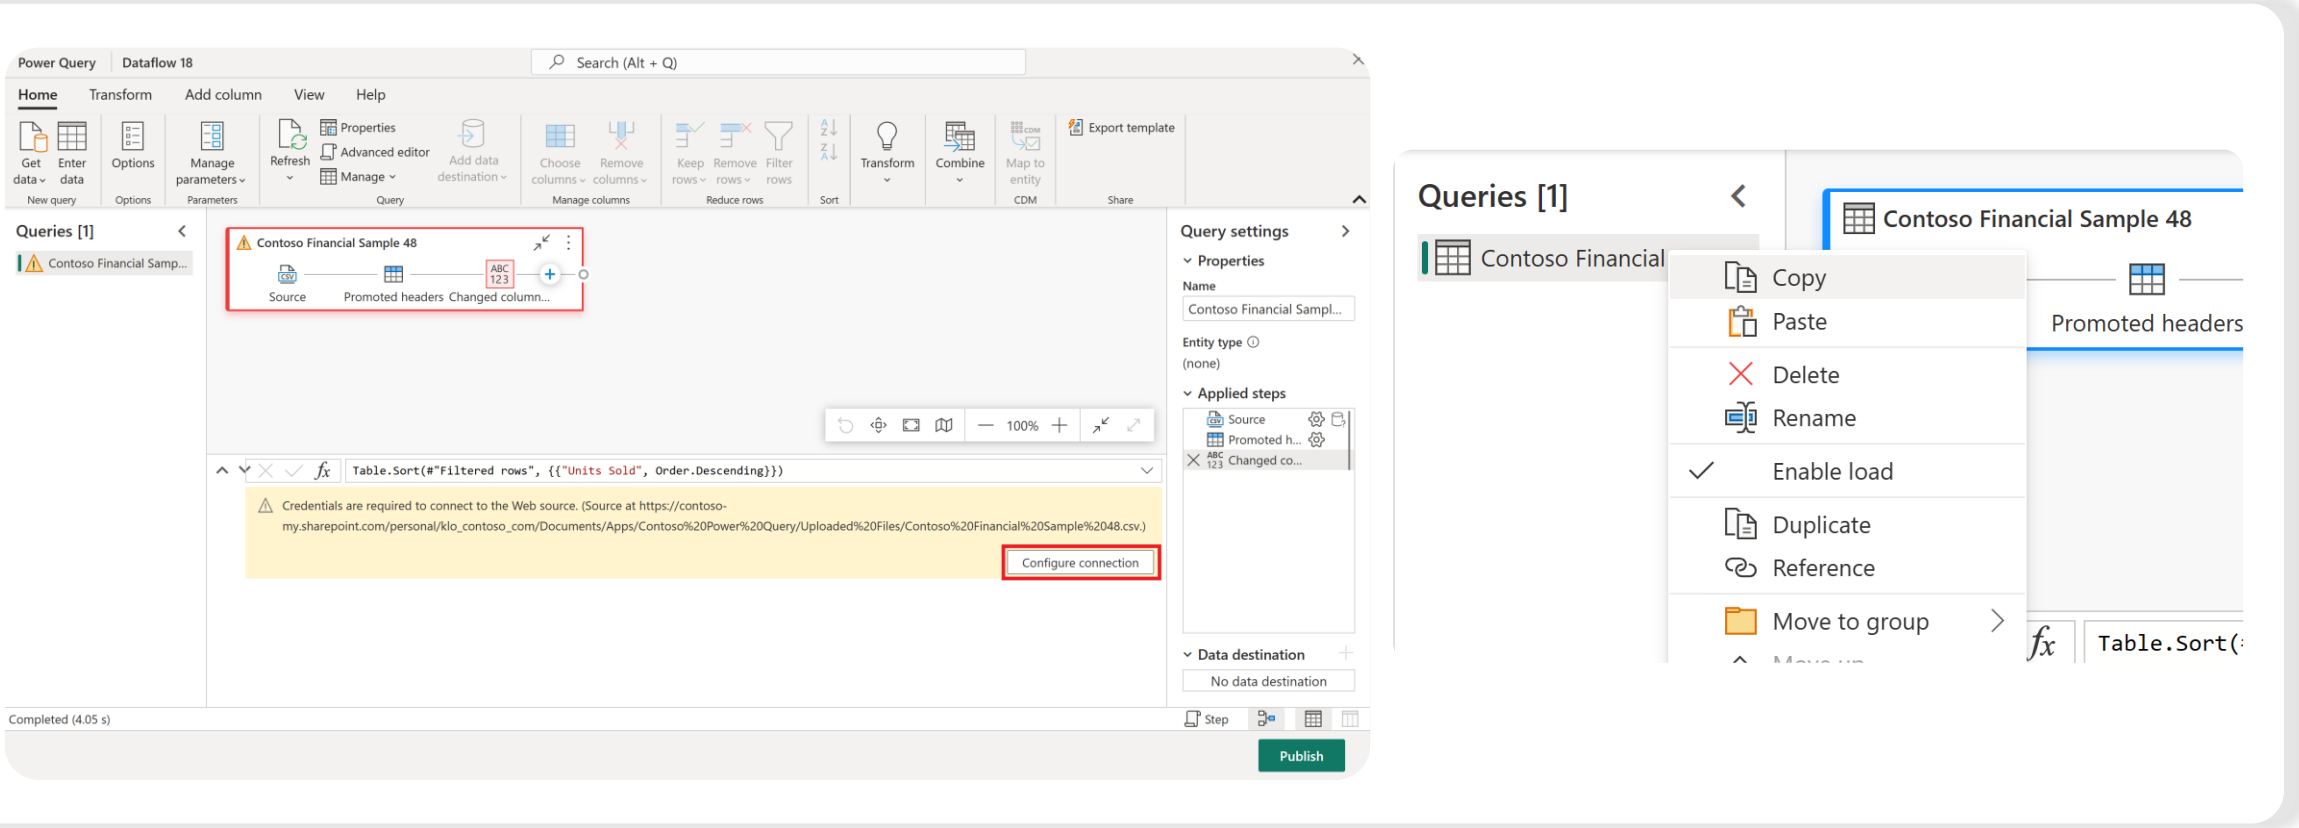

Ingest data into Microsoft Fabric using the Azure Data Factory Copy Activity

Ingest data into Microsoft Fabric using the Azure Data Factory Copy Activity

- Set up authentication
- Write to a Fabric Lakehouse table with an ADF pipeline
- Read from a Fabric Lakehouse table with an ADF pipeline

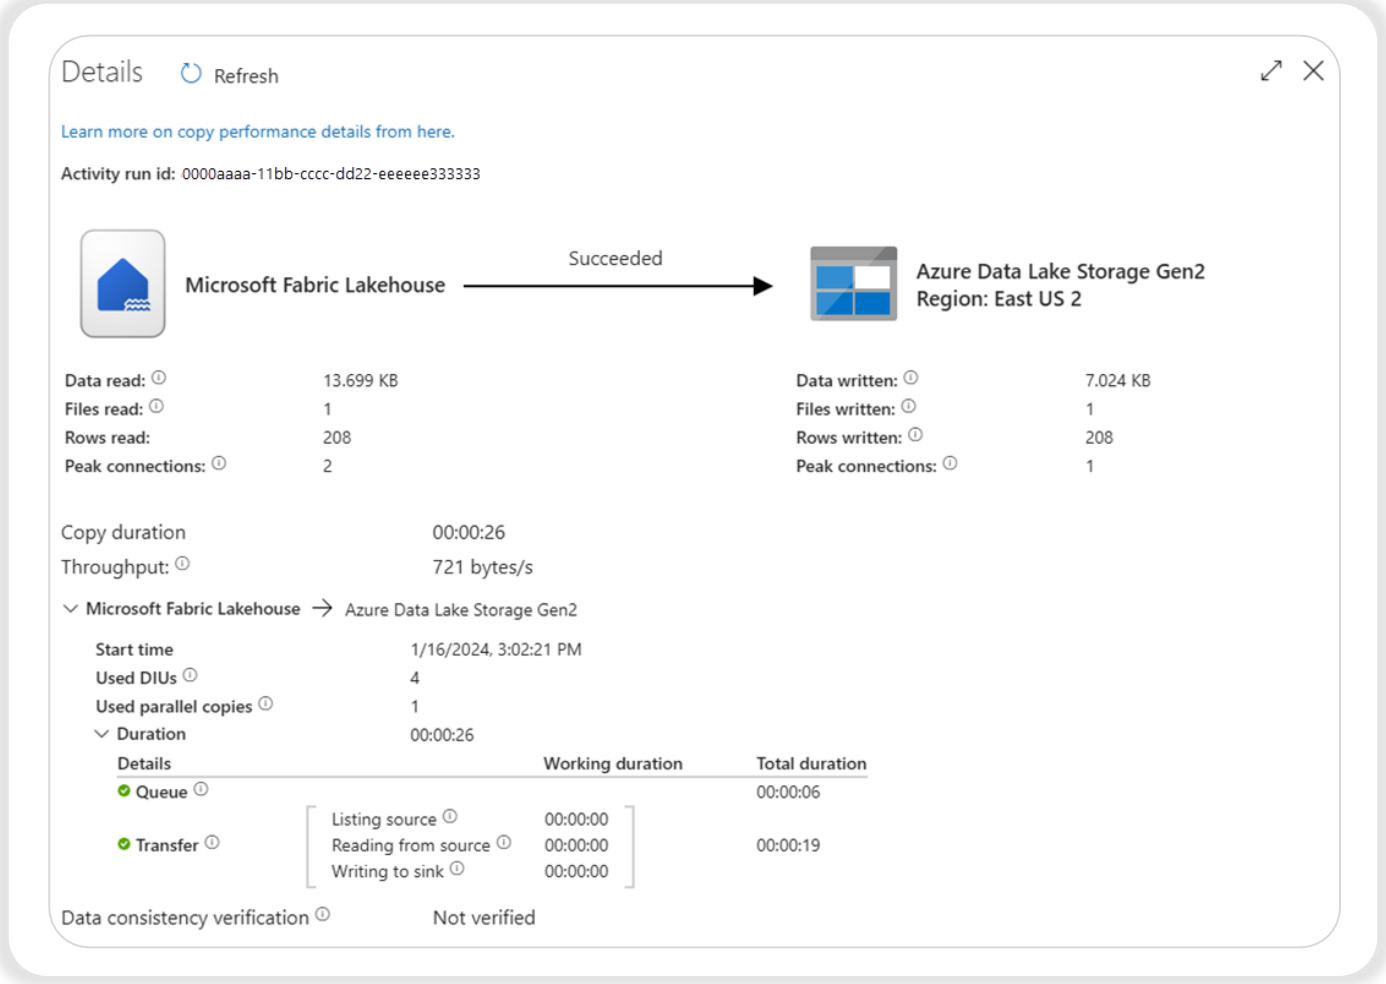

Set up authentication

Create or use an existing app registration service principal (SPN)

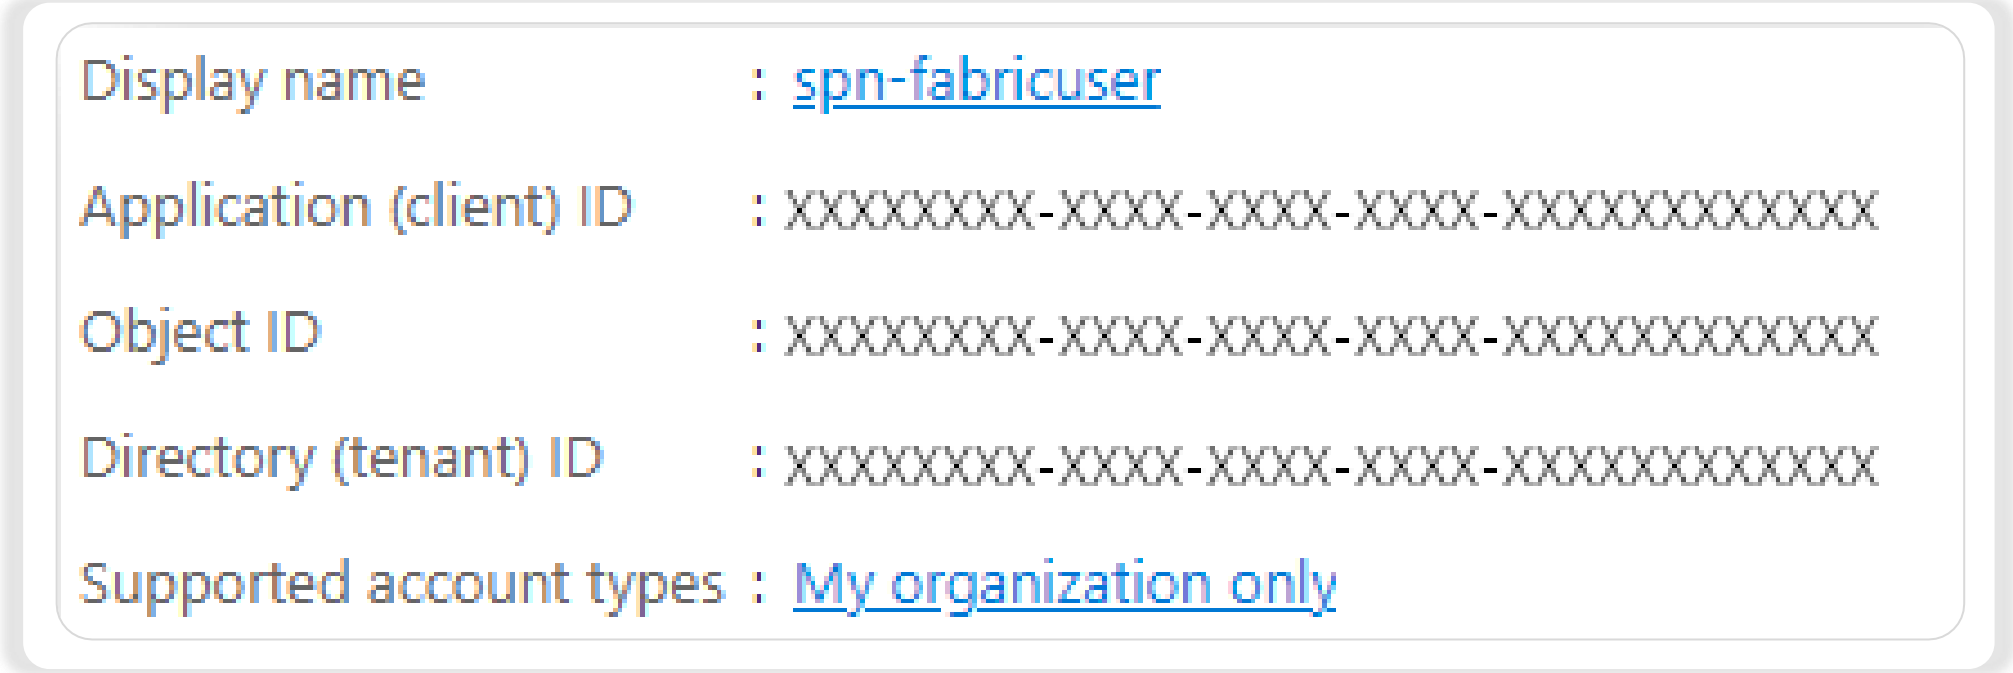

Set up authentication

Create a new Microsoft Entra Security Group or use an existing one then add the SPN to it

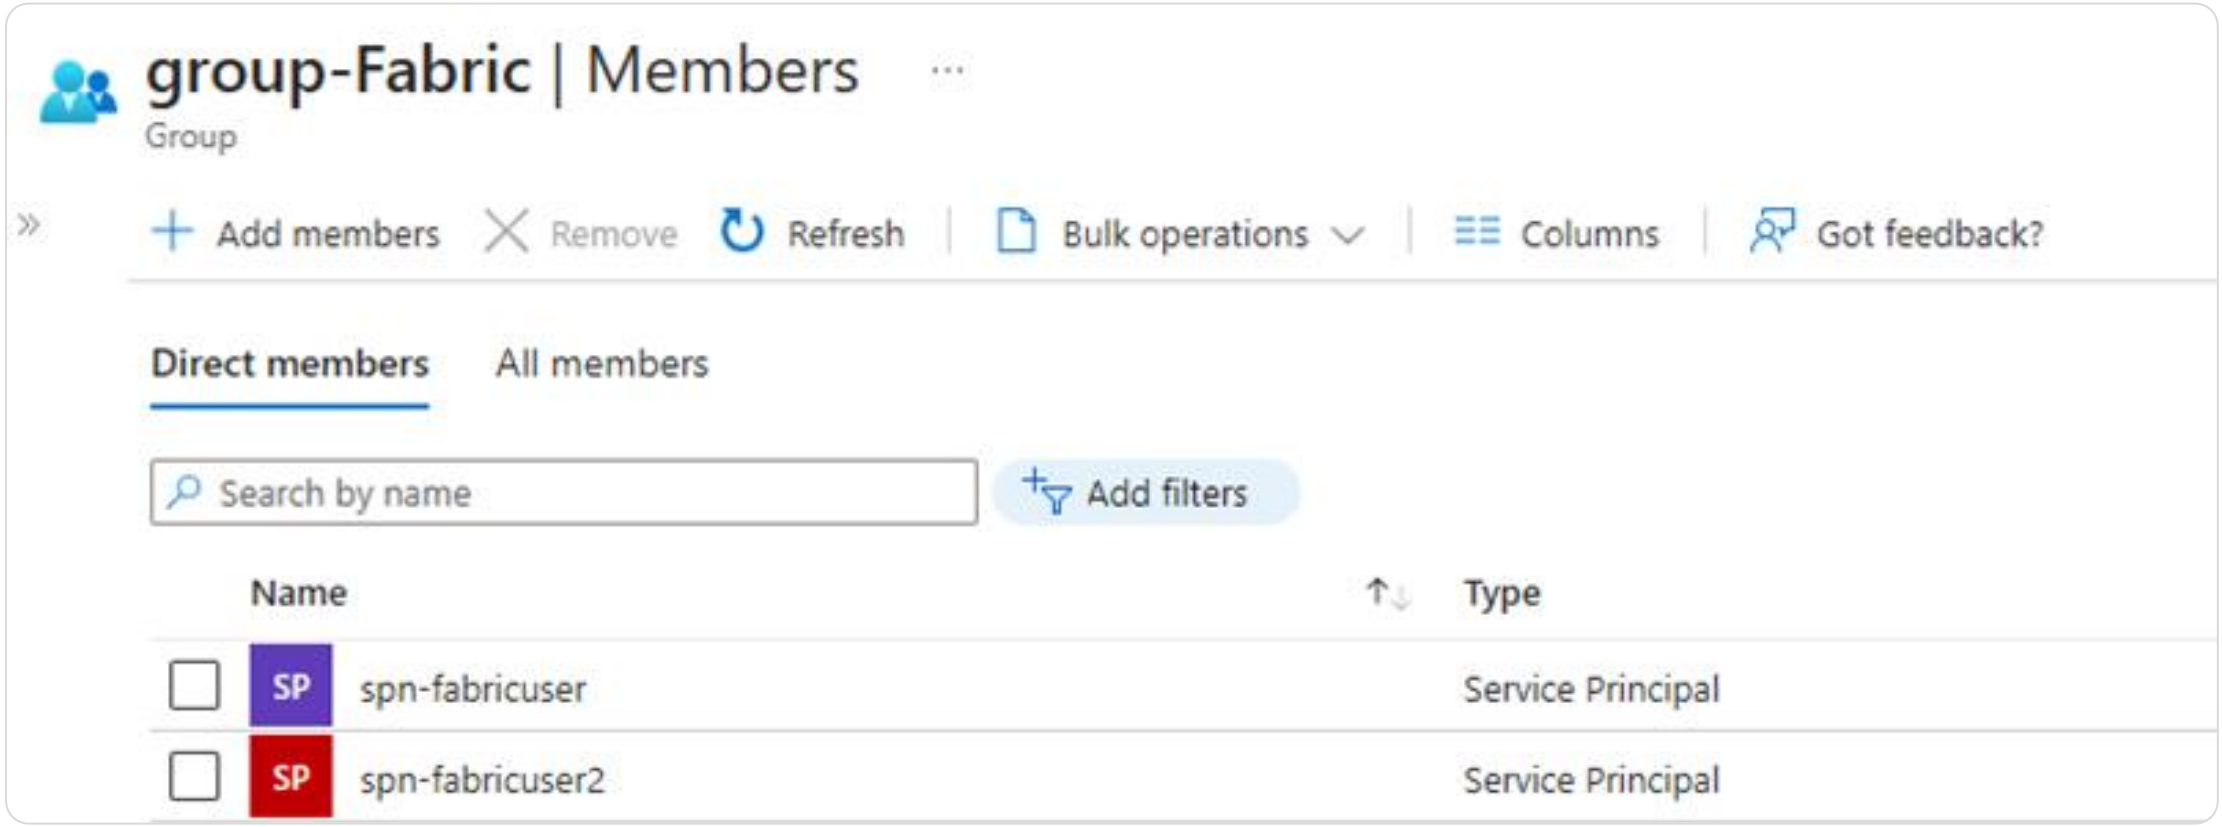

Set up authentication

On Power BI admin portal, **allow service principals and** add the Security Group

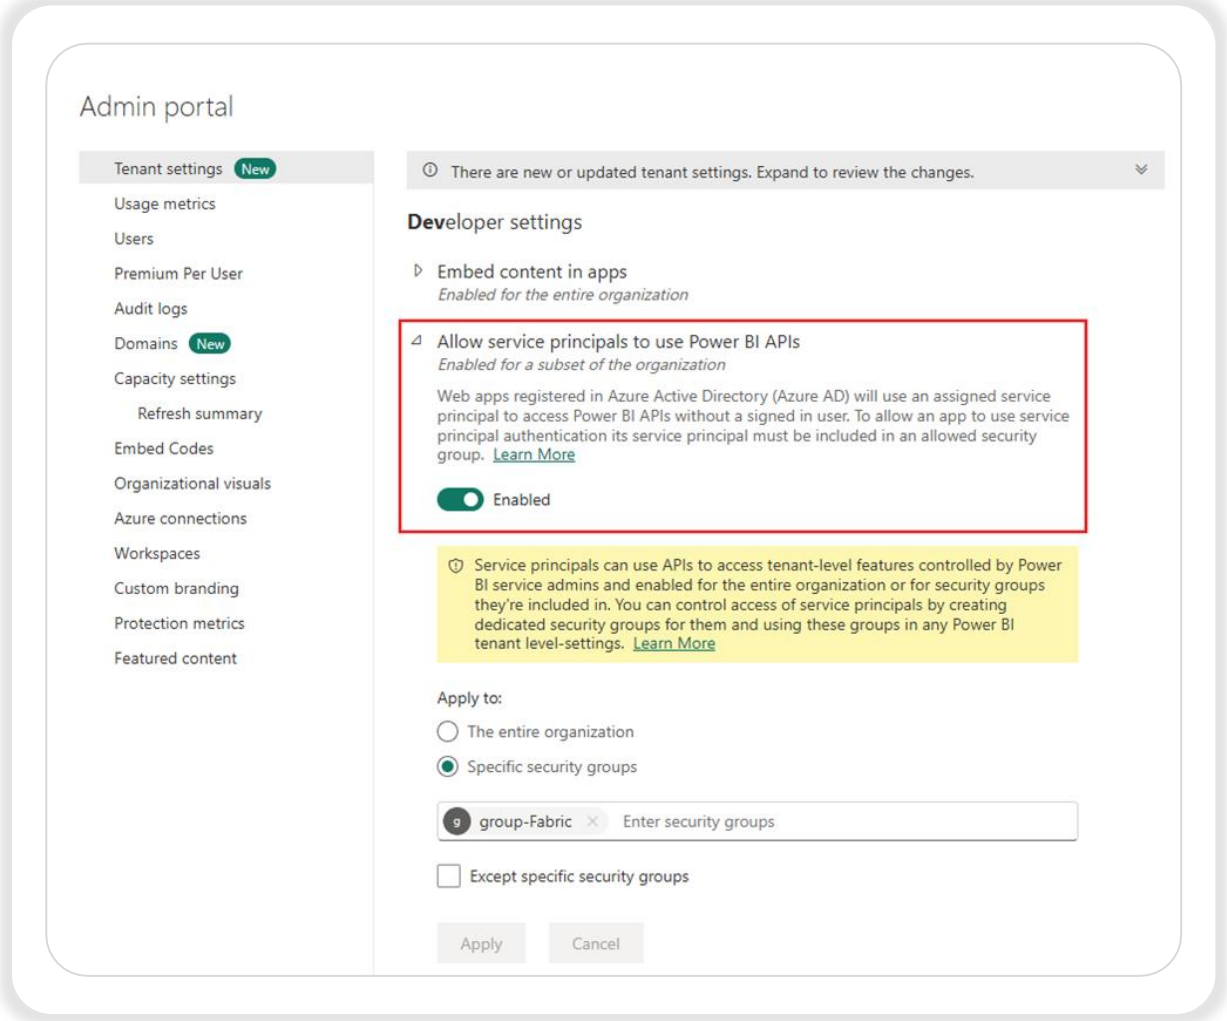

Set up authentication

Add the SPN or the service group to the workspace

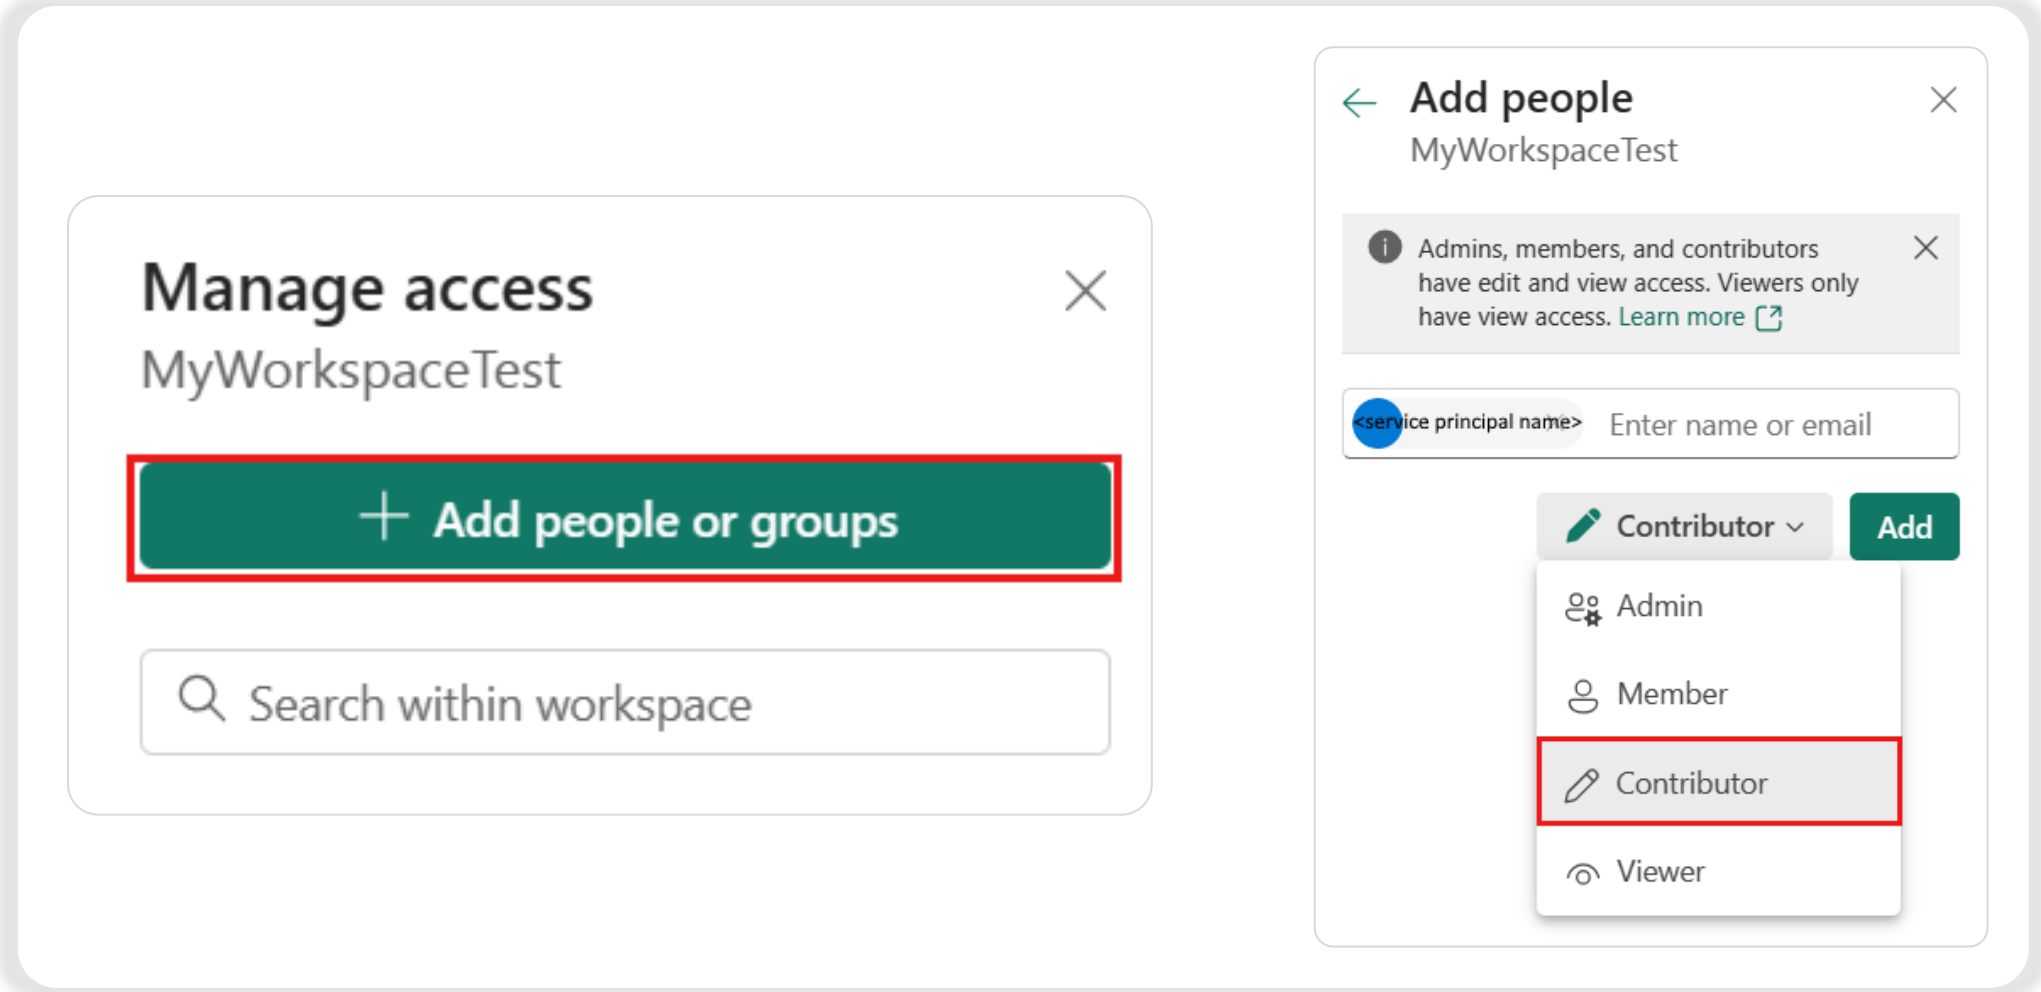

Set up authentication

From Azure Data Factory, create a new Microsoft Fabric Lakehouse linked service

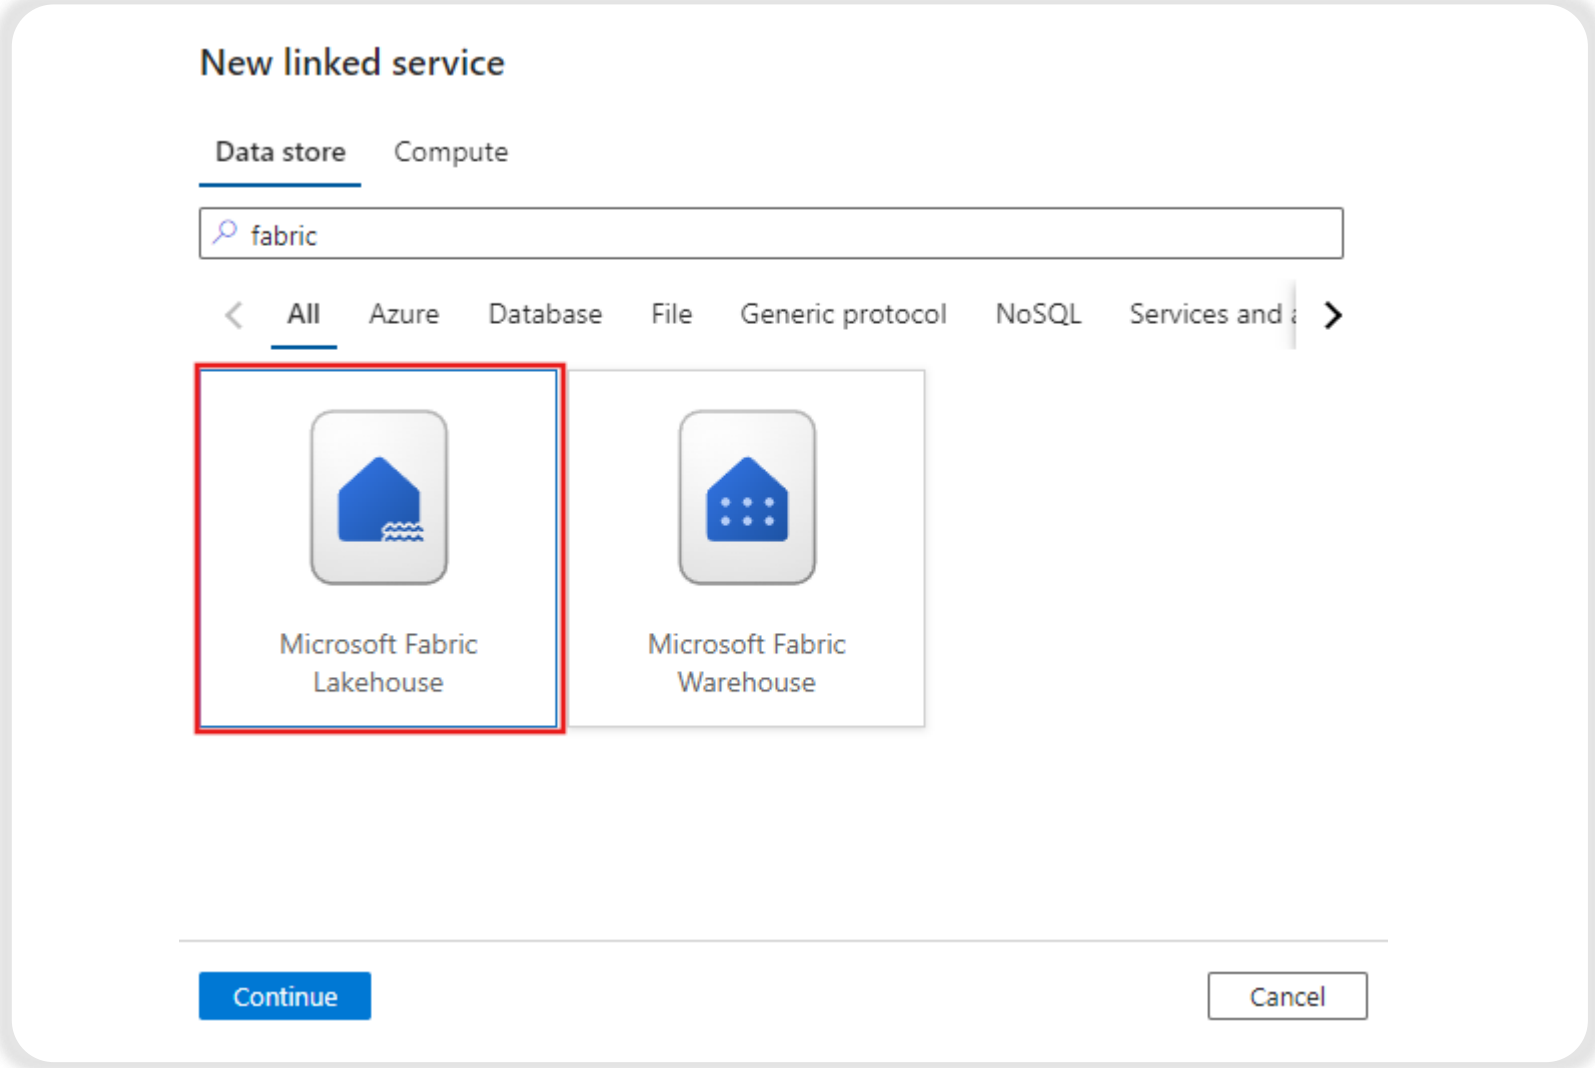

Set up authentication

Create a Dataset that references the Microsoft Fabric Lakehouse linked service

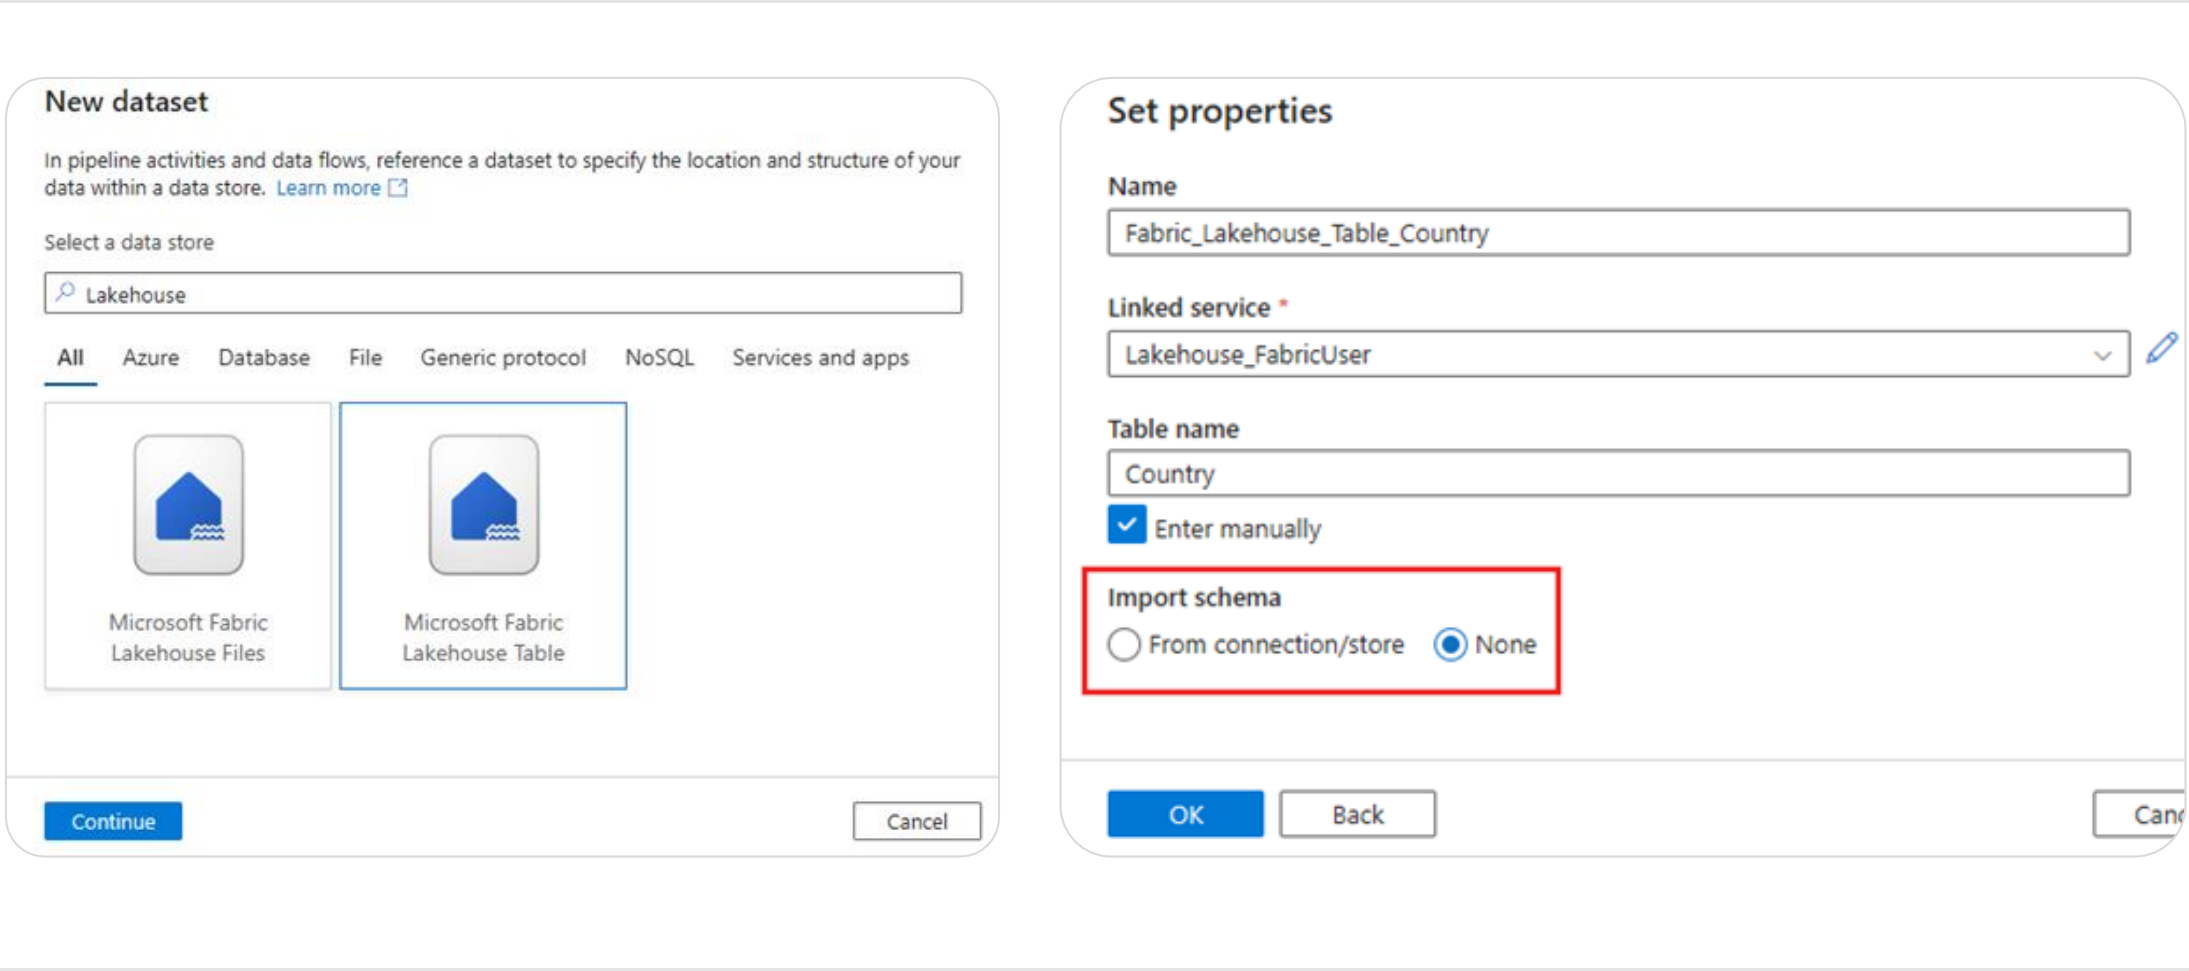

Write to a Fabric Lakehouse table with an ADF pipeline

Create a new pipeline and add a Copy activity to the pipeline canvas

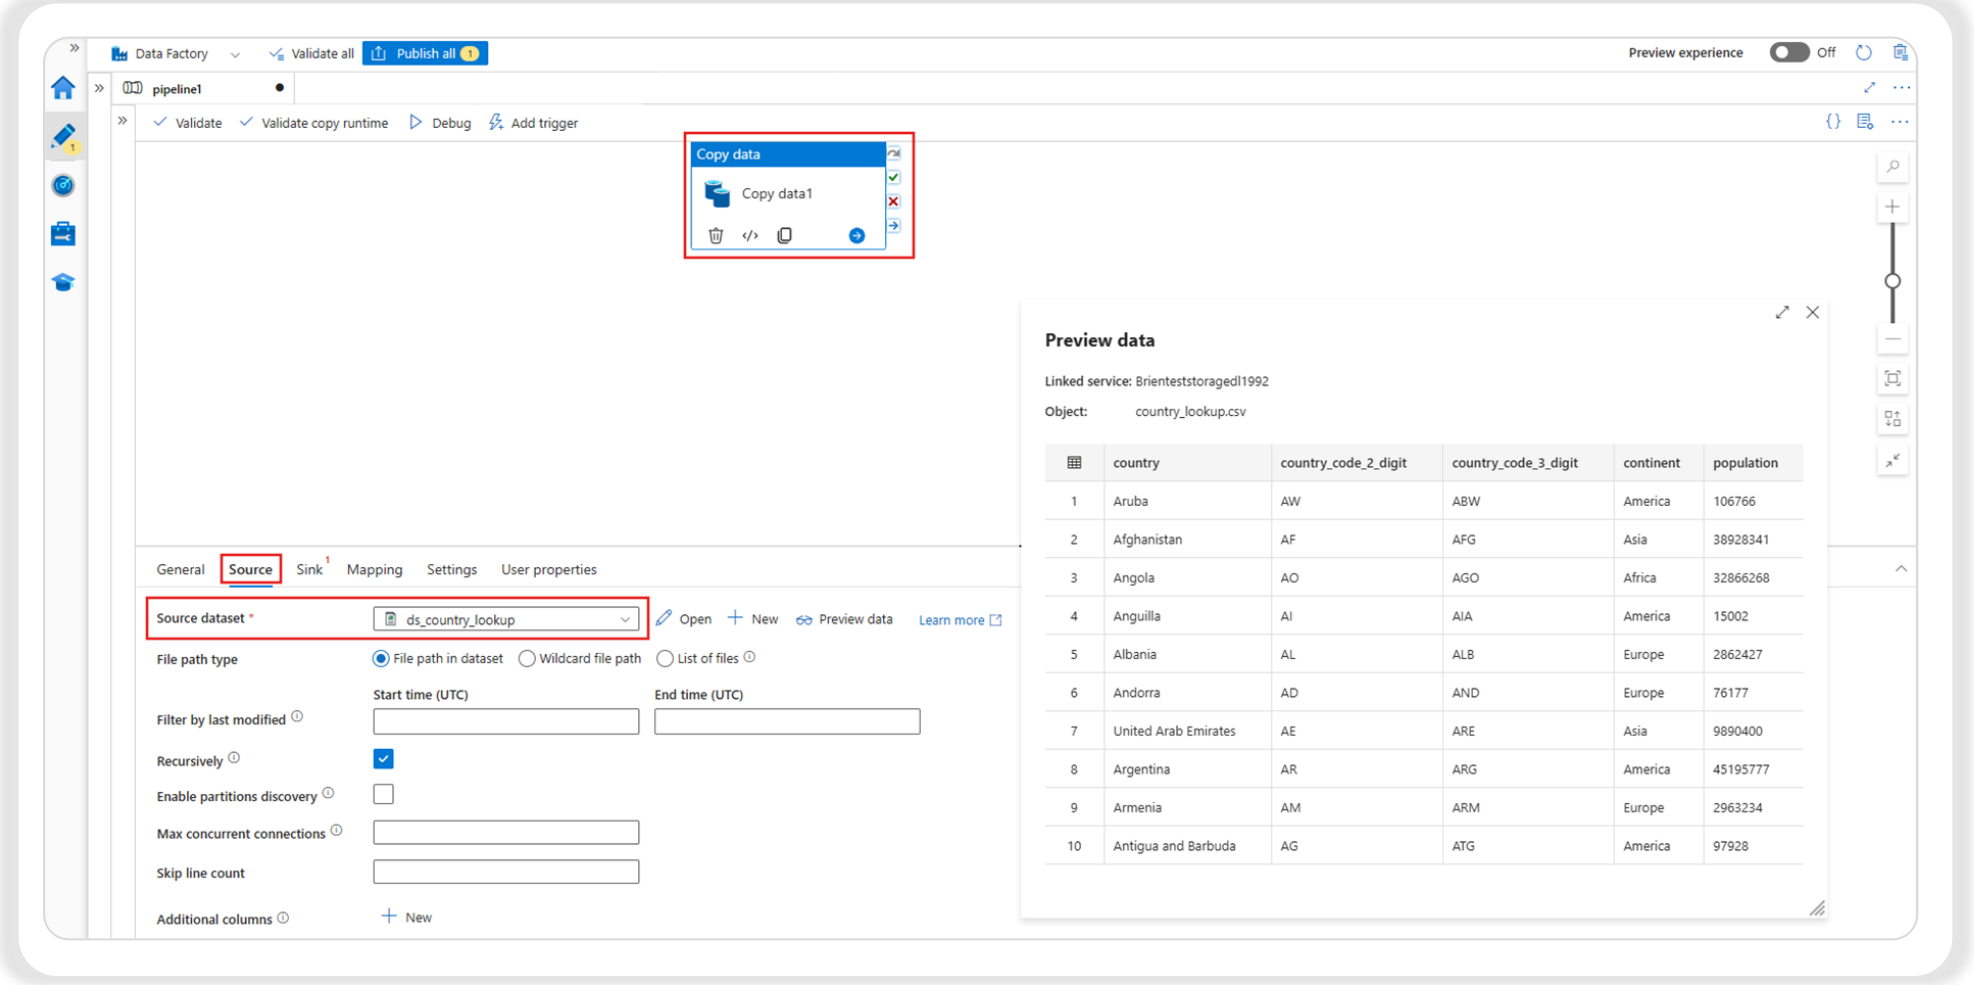

	country	country_code_2_digit	country_code_3_digit	continent	population
1	Aruba	AW	ABW	America	106766
2	Afghanistan	AF	AFG	Asia	38928341
3	Angola	AO	AGO	Africa	32866268
4	Anguilla	AI	AIA	America	15002
5	Albania	AL	ALB	Europe	2862427
6	Andorra	AD	AND	Europe	76177
7	United Arab Emirates	AE	ARE	Asia	9890400
8	Argentina	AR	ARG	America	45195777
9	Armenia	AM	ARM	Europe	2963234
10	Antigua and Barbuda	AG	ATG	America	97928

Write to a Fabric Lakehouse table with an ADF pipeline

From the Copy activity's Sink tab and select the Fabric Lakehouse dataset

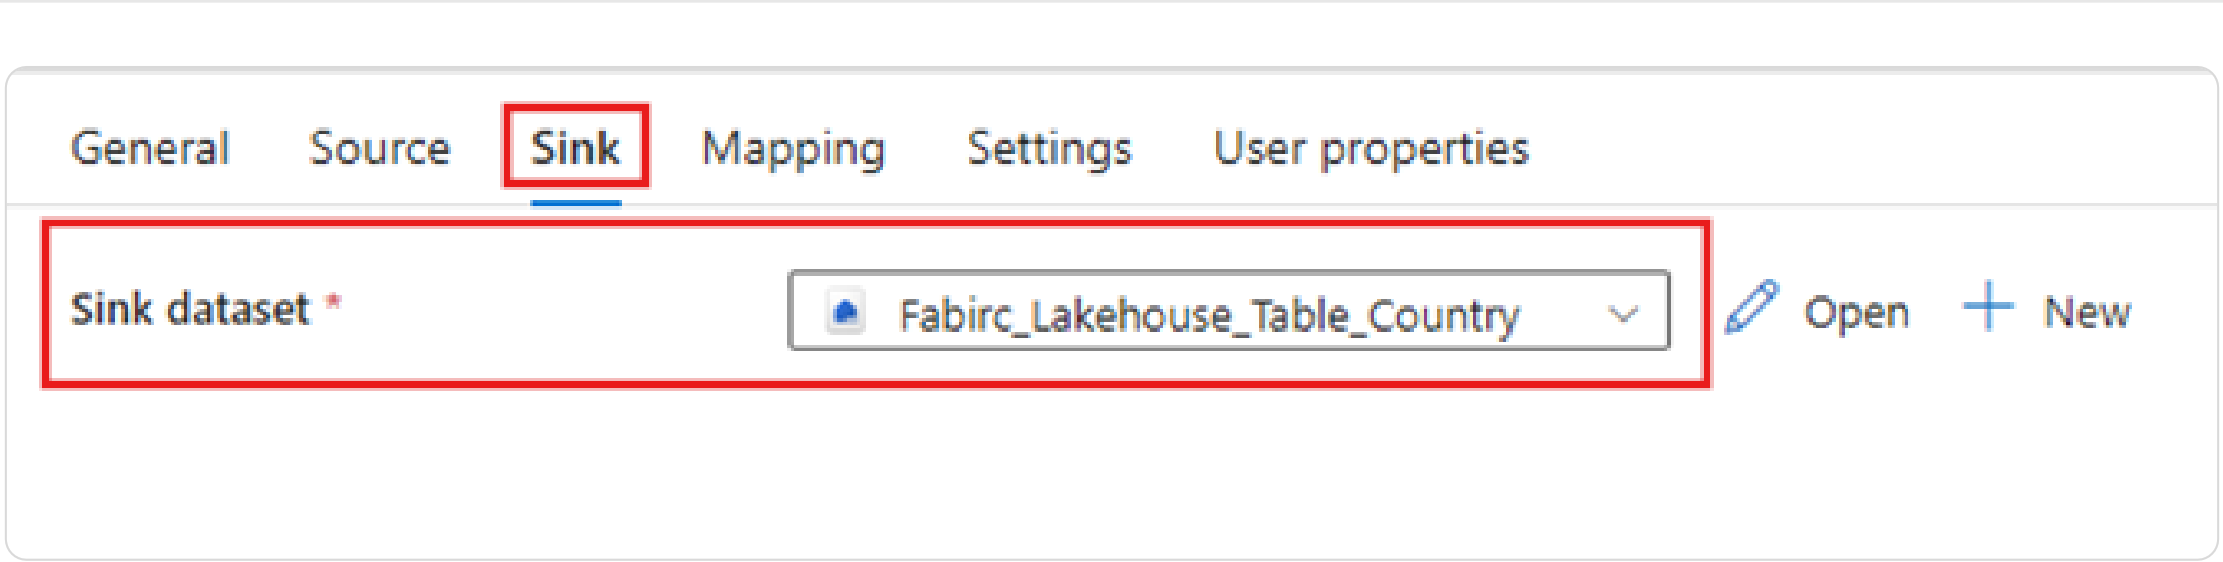

Write to a Fabric Lakehouse table with an ADF pipeline

Run the pipeline

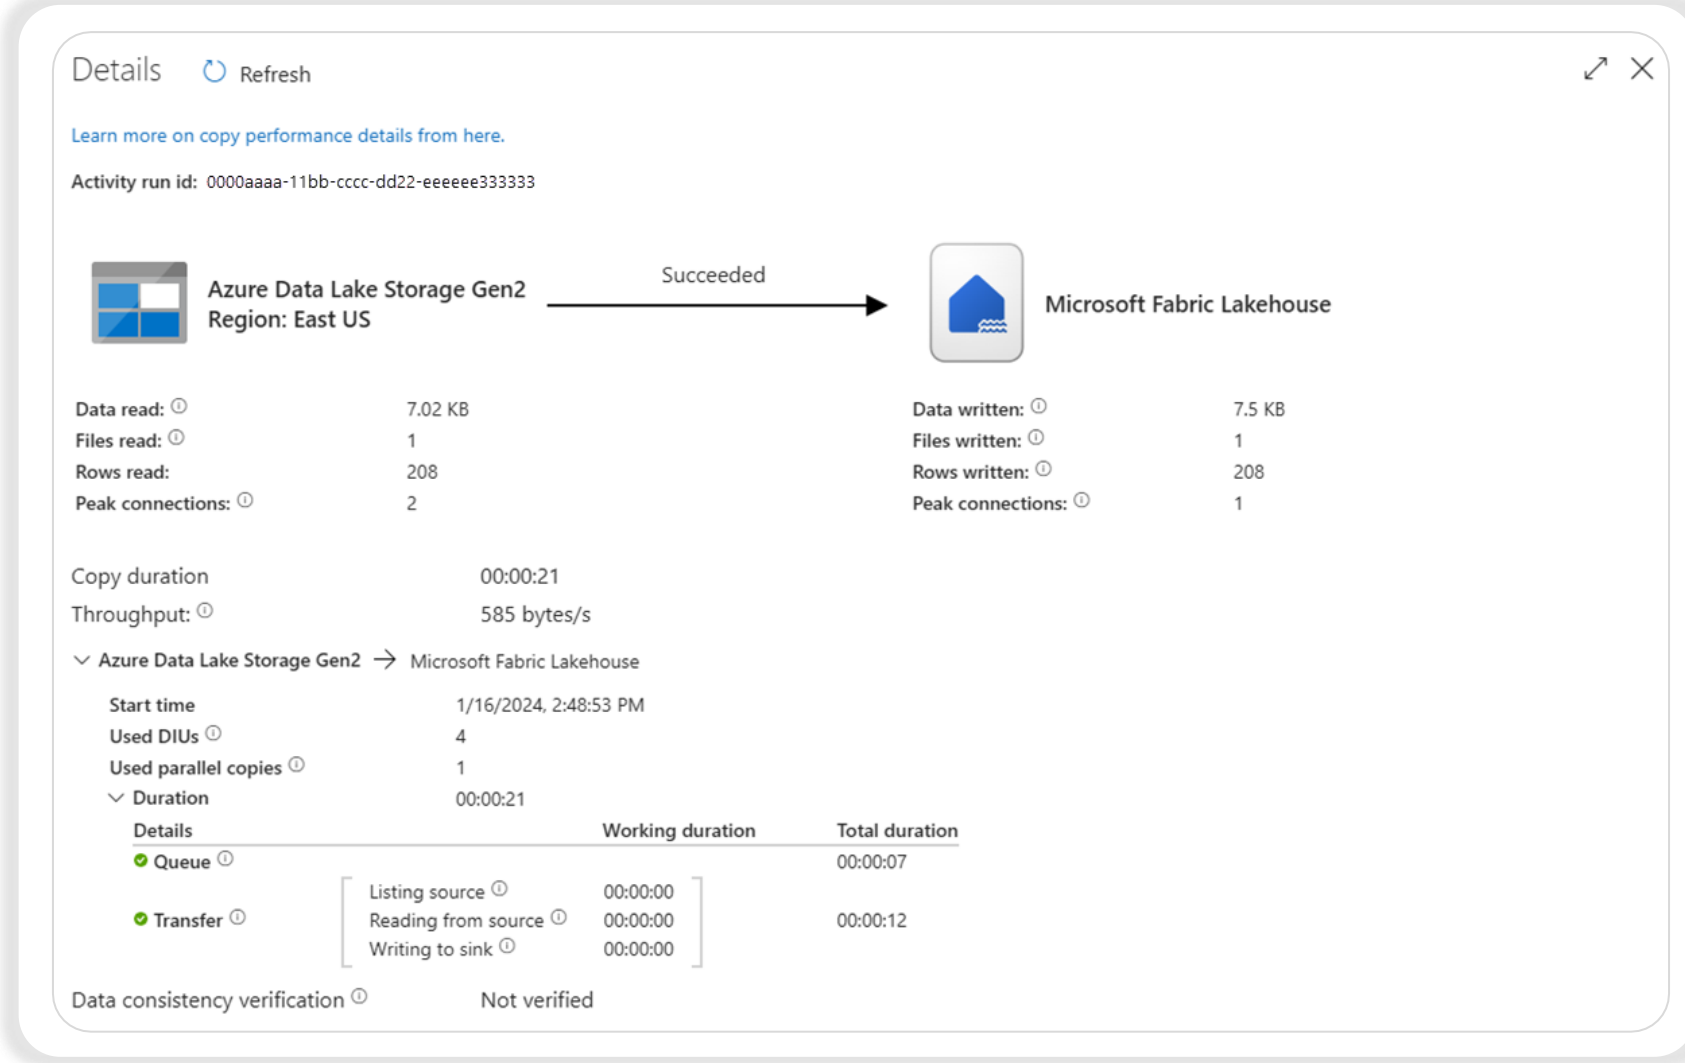

Read from a Fabric Lakehouse table with an ADF pipeline

Create a new pipeline and add a Copy activity to the pipeline canvas

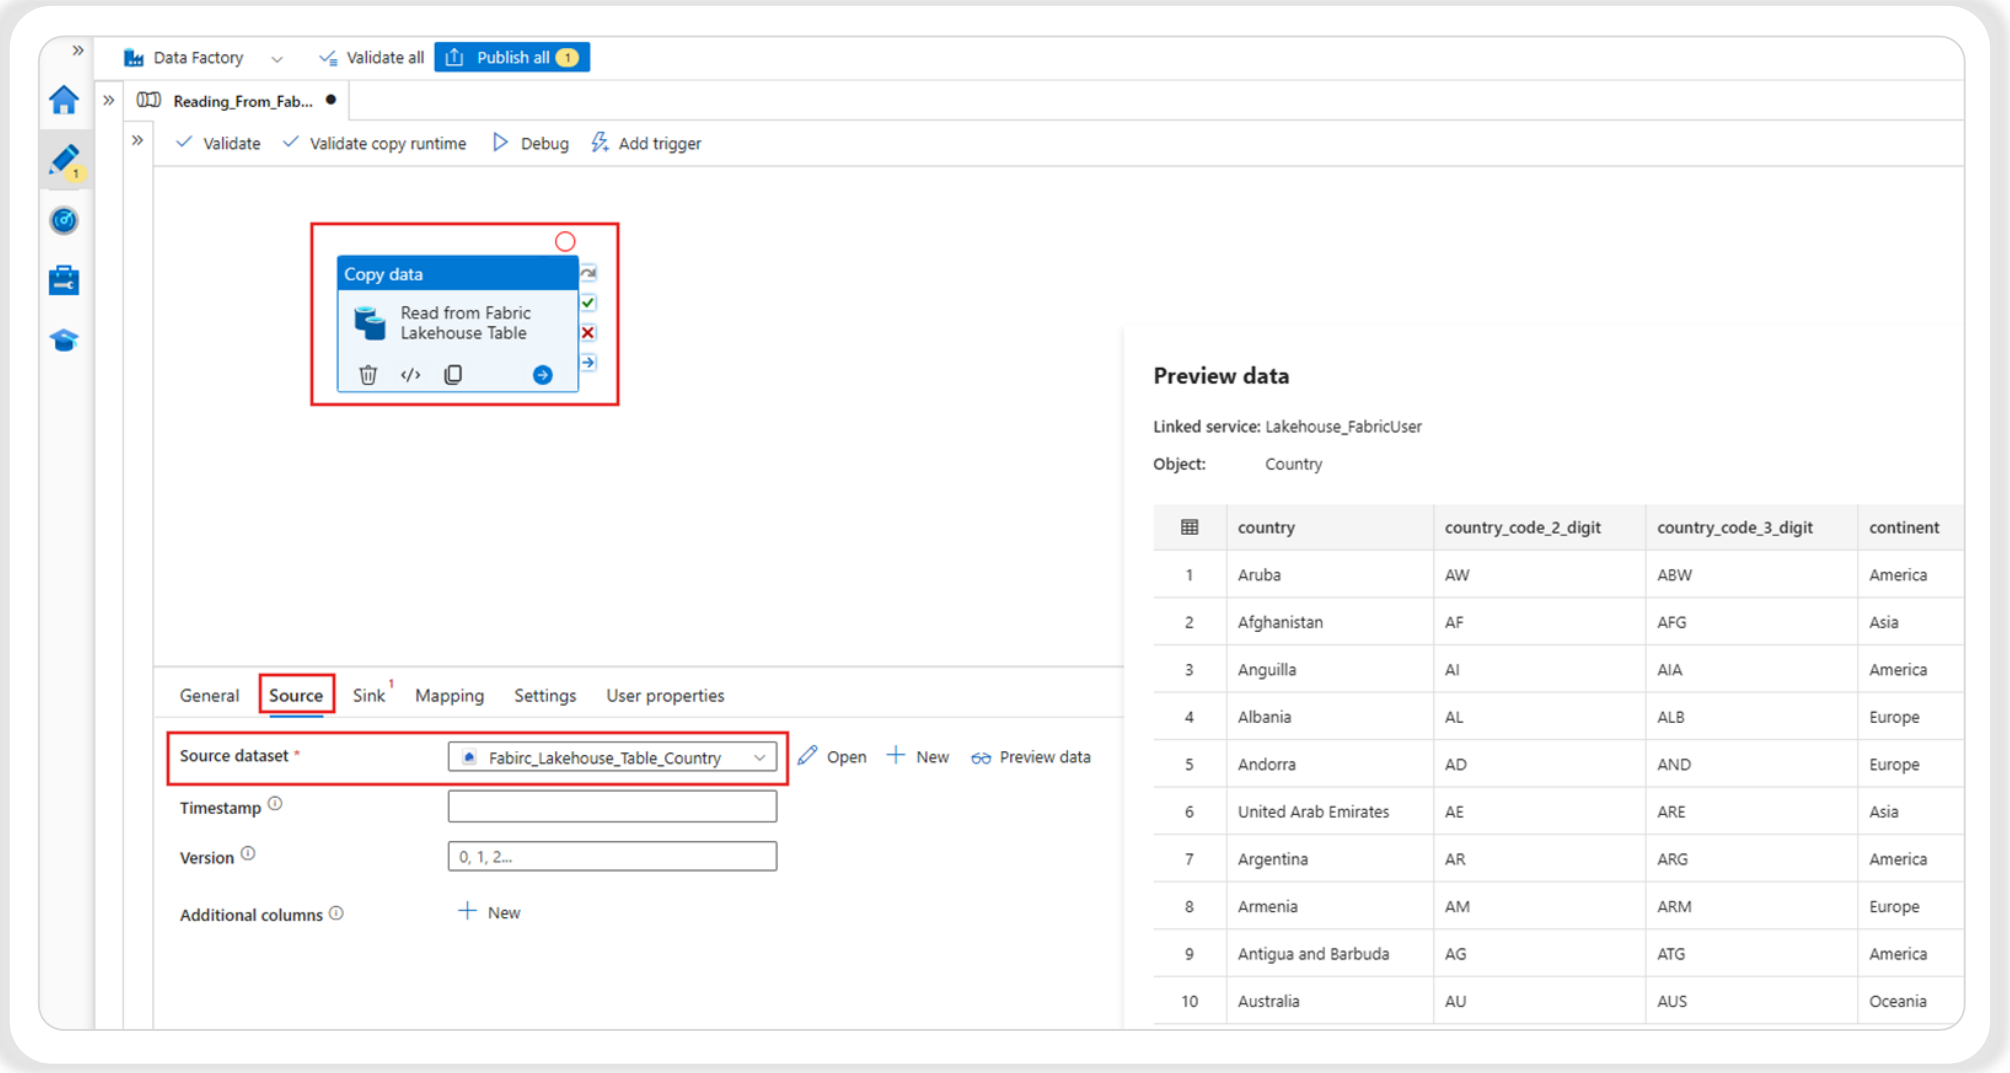

	country	country_code_2_digit	country_code_3_digit	continent
1	Aruba	AW	ABW	America
2	Afghanistan	AF	AFG	Asia
3	Anguilla	AI	AIA	America
4	Albania	AL	ALB	Europe
5	Andorra	AD	AND	Europe
6	United Arab Emirates	AE	ARE	Asia
7	Argentina	AR	ARG	America
8	Armenia	AM	ARM	Europe
9	Antigua and Barbuda	AG	ATG	America
10	Australia	AU	AUS	Oceania

Read from a Fabric Lakehouse table with an ADF pipeline

Copy activity's Sink tab and select the destination dataset

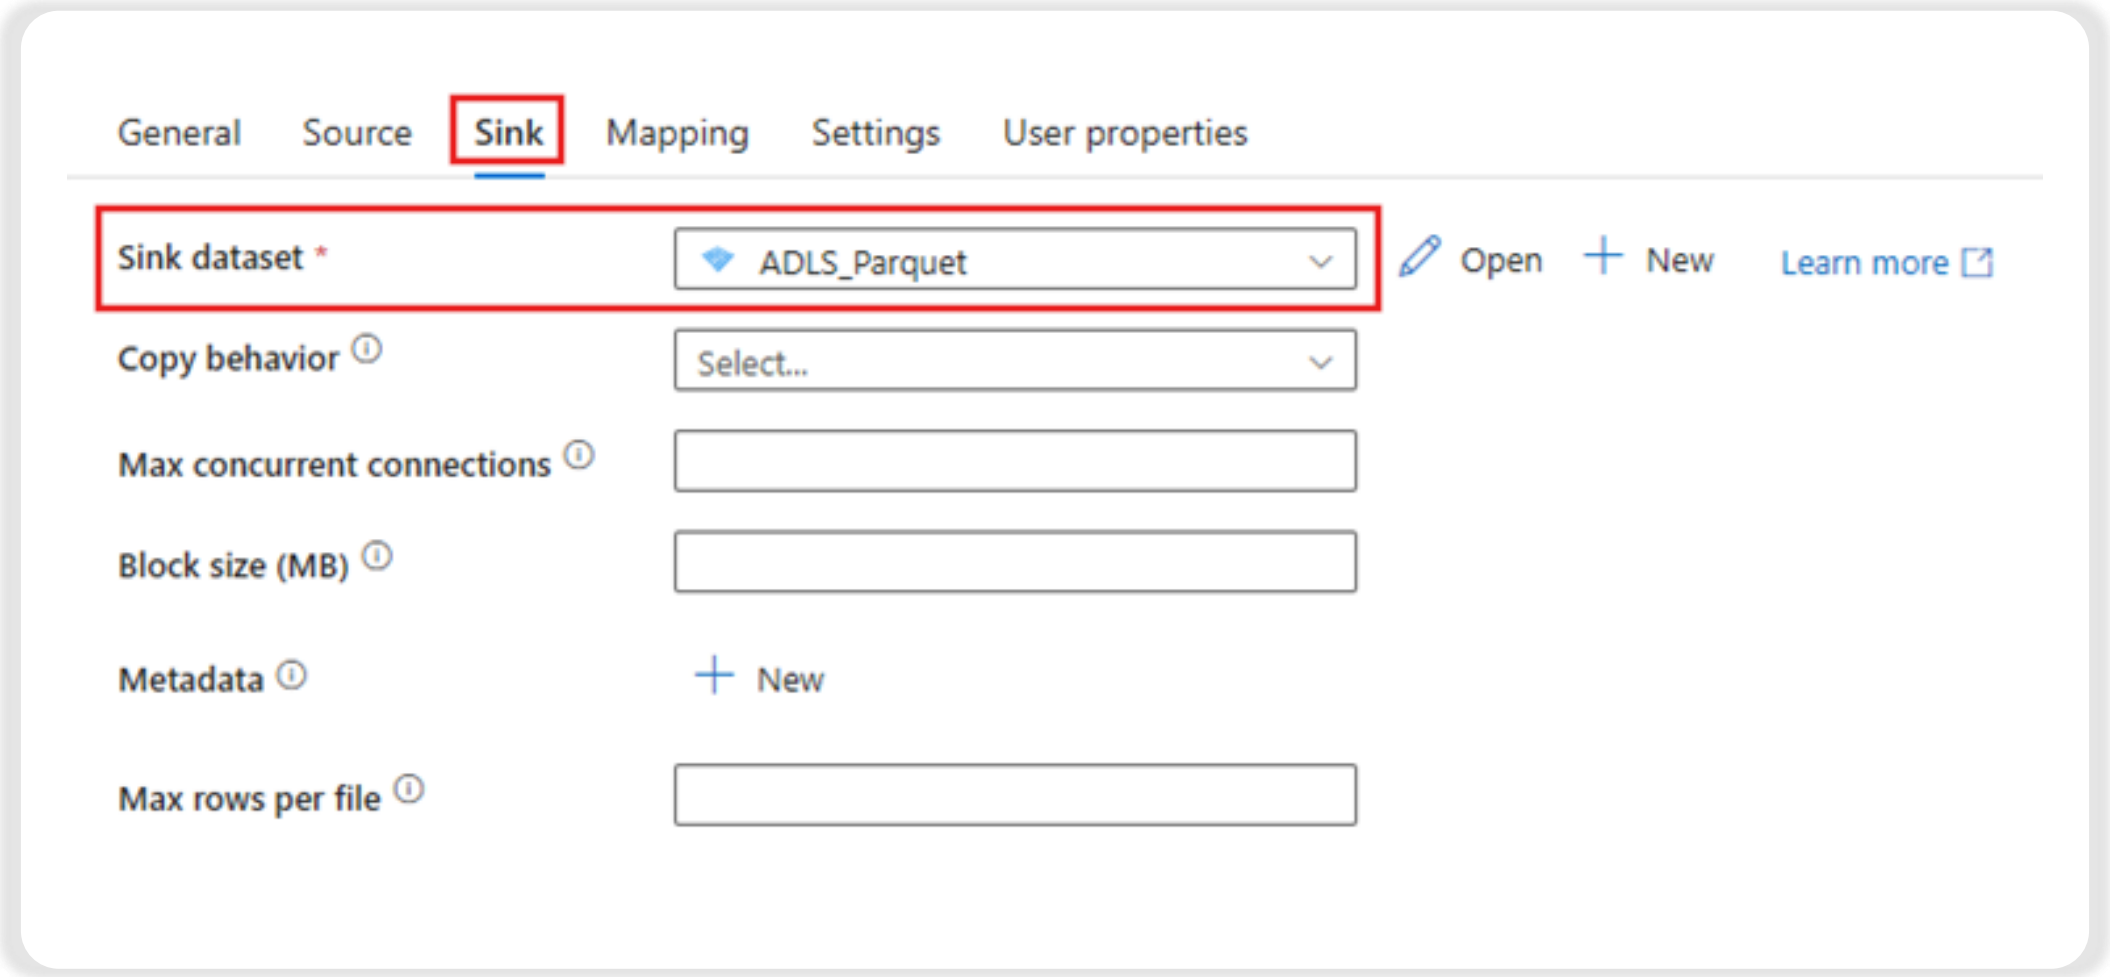

Read from a Fabric Lakehouse table with an ADF pipeline

Run the pipeline

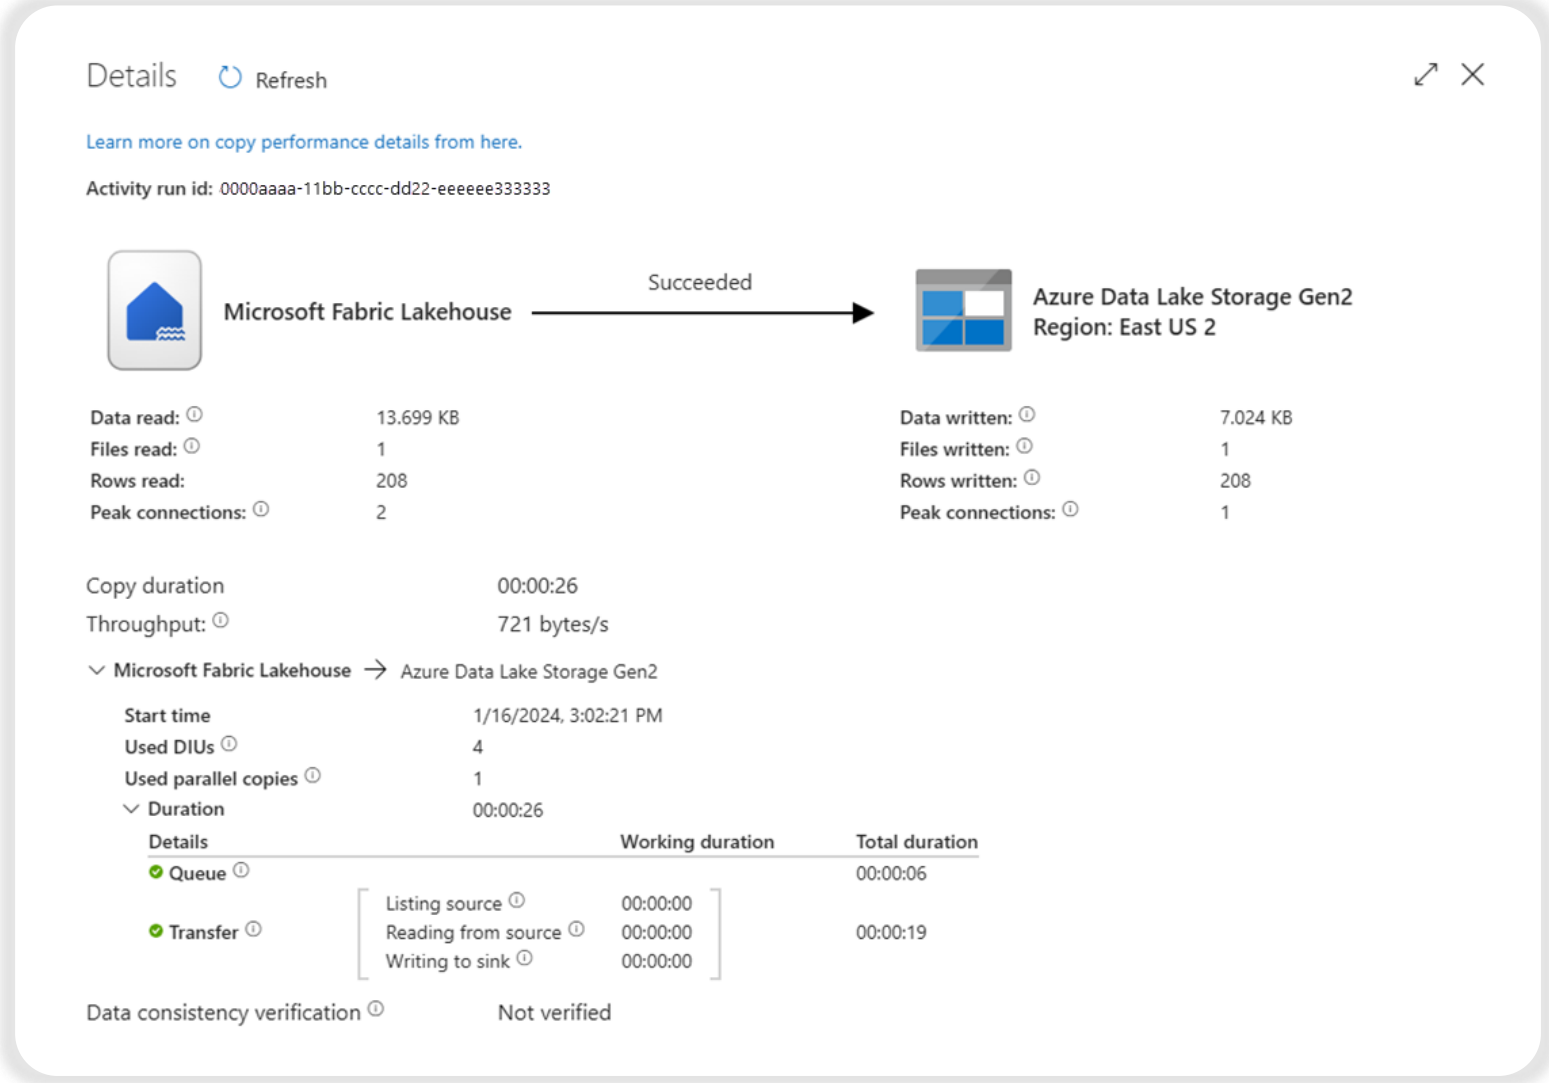

Read from a Fabric Lakehouse table with an ADF pipeline

Inspect the Parquet file in ADLS Gen2

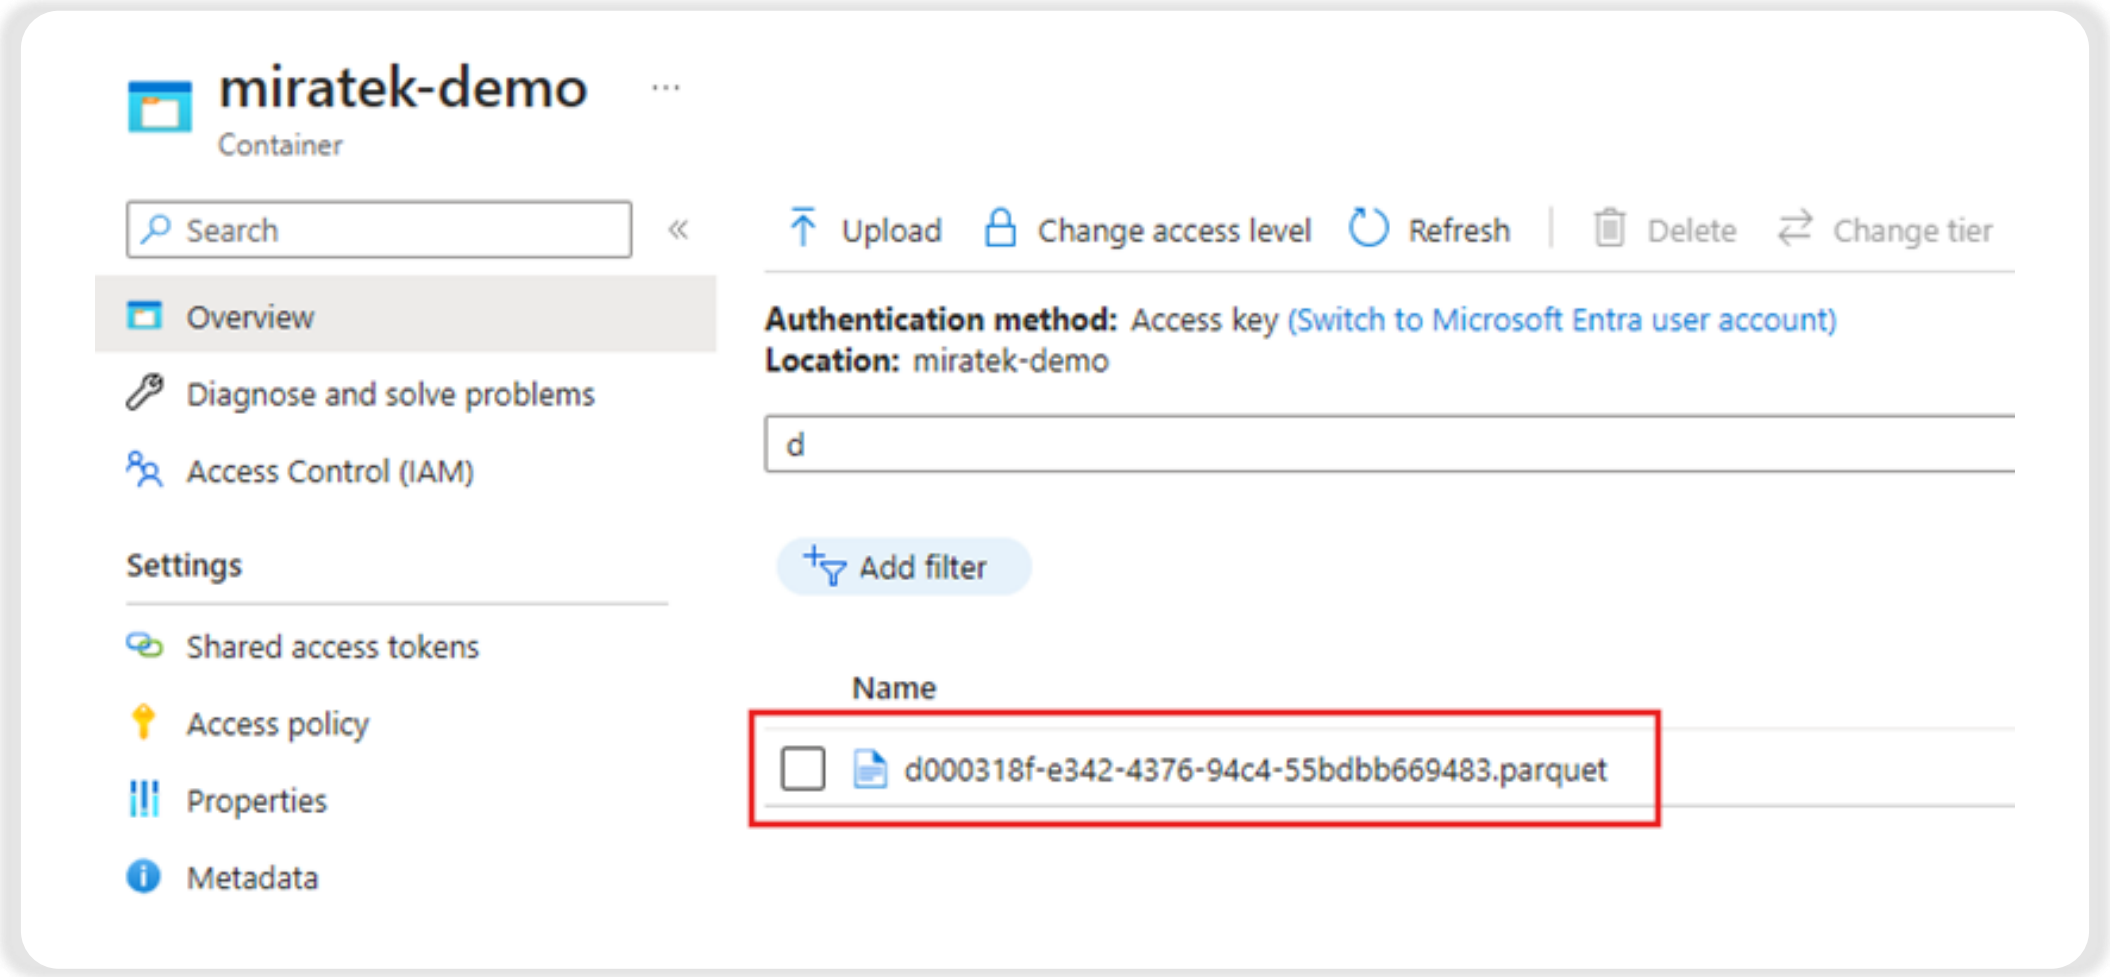

Hands-on labs

Hands-on Lab

Please note that the labs are only open to a limited number of Microsoft partner participants and are offered on a **first-come-first-served basis**.

- **Step 1** : Click on **Launch** in the event curriculum to launch your labs

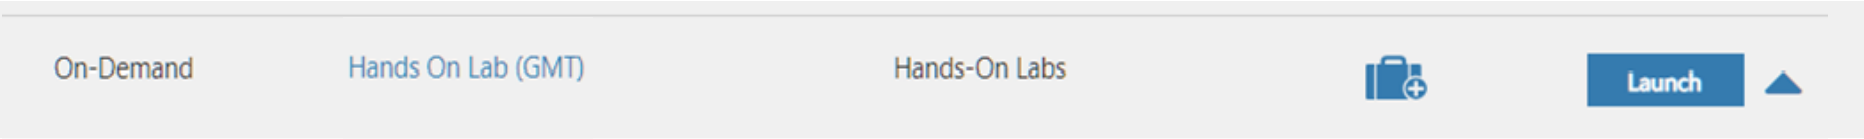

- **Step 2 :** Join the **Live session for Lab Support and Q&A Session**

For all Lab logistical support issues, please go to: https://www.skillable.com/customer-support/ and select the "Open a Support Ticket"

Important lab information

Track remaining lab execution time

- Once launched, a lab is available for active execution for the time displayed in the screenshot.

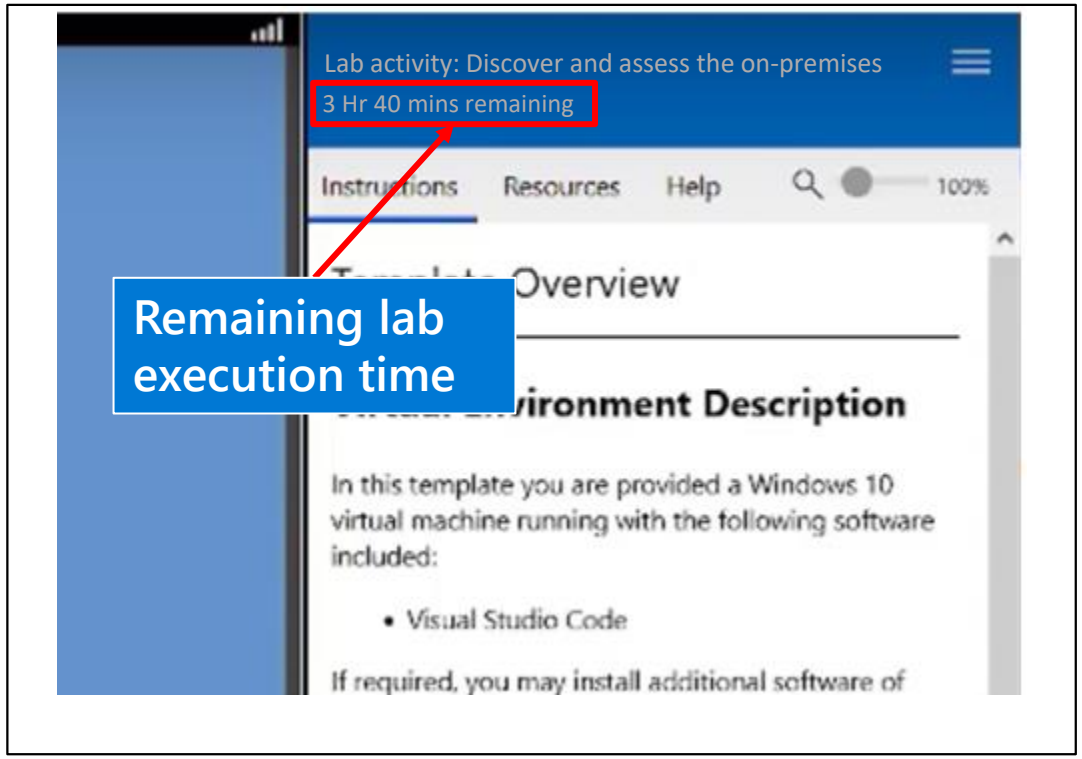

Save your Lab

- Ensure that you SAVE the lab after each break.
- Do not fail to relaunch within the specified time to avoid losing saved labs.

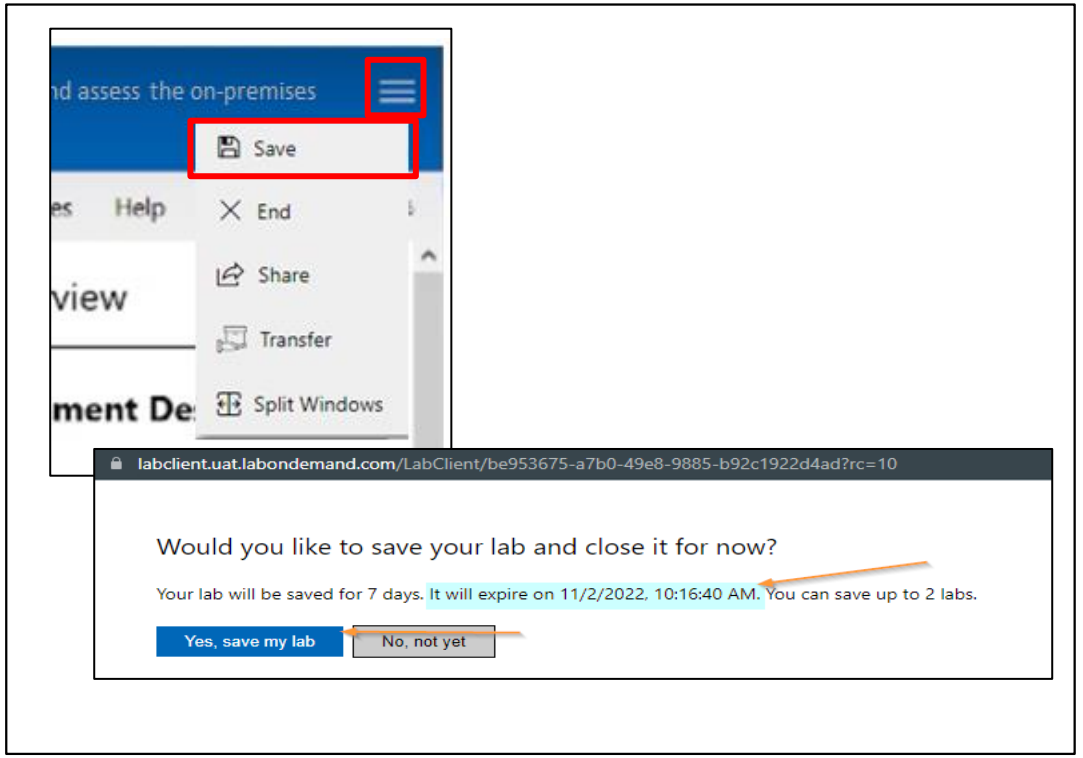

Once the allocated lab time expires, it can not be extended.

Hands-on labs

- **Lab 02**

  Use Case 02: Data Factory solution for moving and transforming data with dataflows and data pipelines

Microsoft

Technofocus

Thank you